

A framework for automatic IT architecture modeling: applying truth discovery

Margus Välja · Robert Lagerström ·

Ulrik Franke · Göran Ericsson

Received: date / Accepted: date

Abstract Modeling IT architecture is a complex, time consuming, and error prone task. However, many systems produce information that can be used in order to automate modeling. Early studies show that this is a feasible approach if we can overcome certain obstacles. Often more than one source is needed in order to cover the data requirements of an IT architecture model and the use of multiple sources means that heterogeneous data needs to be merged. Moreover, the same collection of data might be useful for creating more than one kind of model for decision support.

IT architecture is constantly changing and data sources provide information that can deviate from reality to some degree. There can be problems with varying accuracy (e.g. actuality and coverage), representation (e.g. data syntax and file format), or inconsistent semantics. Thus, integration of heterogeneous data from different sources needs to handle data quality problems of the sources. This can be done by using probabilistic models. In the field of truth discovery, these models have been developed to track data source trustworthiness in order to help solving conflicts while making quality issues manageable for automatic modeling.

Margus Välja, Robert Lagerström, Göran Ericsson
KTH Royal Institute of Technology
100 44 Stockholm, Sweden
E-mail: {margusv, robertl, ergo}@kth.se

Ulrik Franke
RISE SICS – Swedish Institute of Computer Science
164 40 Kista, Sweden
E-mail: ulrik.franke@ri.se

We build upon previous research in modeling automation and propose a framework for merging data from multiple sources with a truth discovery algorithm to create multiple IT architecture models. The usefulness of the proposed framework is demonstrated in a study where models using three tools are created, namely; Archi, securiCAD, and EMFTA.

Keywords IT architecture modeling · System modeling · Automatic data collection · Automatic modeling

1 Introduction

Modern enterprise architecture is complex [25]. Systems and software that have a prevalent role in everyday tasks are interconnected and used in different ways by employees, partners, and suppliers. Moreover, these software and systems change when an organization keeps evolving. Models, representations that explain and characterize real system entities, can be employed to understand and manage such a complex reality. However, in order to base important decisions on models, the models need to be created using updated information [17]. For large organizations with many interconnected systems and software this is difficult, time consuming, and error prone.

Currently, many available approaches for enterprise modeling are manual. At the same time there is a vast amount of data available within organizations that could be used for modeling. The availability of this data makes automation of modeling possible, and helps to capture and evaluate dynamic changes in enterprise IT architecture with less effort. There have been attempts to collect data for automatic modeling in previous studies. Buschle et al. [5] use data collected with a vulnerability scanner to instantiate a metamodel for IT architecture security evaluation. In another study, Holm et al. [21] utilize a similar method to generate an ArchiMate model. The authors in the two studies show that the proposed method is applicable and can be used to generate sufficiently accurate models much faster than with manual modeling. However, the authors in both papers conclude that while the data sources they used were providing good enough data, only a subset of all the metamodel entities could be instantiated. To create more complete models additional data sources are needed. The need for multiple data sources for enterprise system modeling has been mentioned by multiple authors such as [34, 9, 46, 36]. Thus, modeling au-

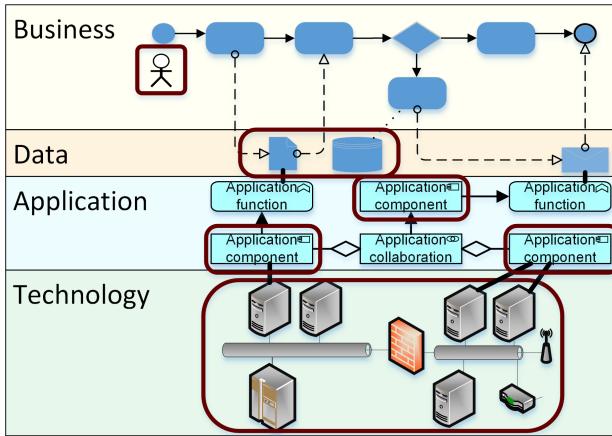


Fig. 1 Mainly the technology and partly the application enterprise layers are studied in the paper. The studied elements are highlighted with dark red rectangles.

tomation would benefit from using multiple data sources. However, the sources often provide data with different scope and representation.

Niemi and Pekkola [35] state that artifacts like architecture models need to cover multiple domains. According to Frank [12], an enterprise can be represented with a set of domain specific interrelated diagrams that represent domain models and are connected through shared concepts. Various enterprise models are described in [43]. A logical conclusion is that an automated modeling approach would not only benefit from using multiple data sources, but also from being able to support the creation of more than one model representing multiple enterprise domains. In our work we focus on IT architecture modeling that covers application components, data, users and IT infrastructure on a holistic enterprise level. The main reason for this delimitation is the availability of data that easily can be automated. However, we believe that the framework supports inclusion of other layers of enterprise architecture as well. Fig. 1 shows the scope of the enterprise modeling covered in this paper using an ArchiMate standard¹ inspired figure. This leads to our first research question: *How can automated system modelling use multiple and heterogeneous sources in an extensible way?*

To ensure successful data integration, different data sets must have, or be given, a common semantic and structural representation [39], [38]. Data needs to be transformed to the same granularity level (the same level of abstraction)

¹ <http://pubs.opengroup.org/architecture/archimate2-doc>

[9] and the conflicts between different data sets need to be resolved. Sources from different domains, however, provide data about various aspects of an enterprise, with varying accuracy (e.g. actuality and coverage) and representation (e.g. data syntax and file format). Data semantics can be inconsistent from one data source to another. Thus, data sets from different sources can be attributed varying degrees of trustworthiness. To manage the quality of data this trustworthiness should be taken into account when an unknown number of heterogeneous data sources are used for automated model creation. There is much literature that deals with data integration e.g. [39,38,8] and trust in data sources such as [37,49,40,29,3,30]. While some of the techniques for data source trustworthiness management described by the authors could be extended to the enterprise domain, none of the authors have investigated this explicitly, leading to our second research question: *How can trustworthiness of multiple and heterogeneous sources be included in automated system modelling?*

The authors have not found a solution in the existing literature that addresses the research questions in an integrated manner. Thus, there is a gap in the field, which we aim to address by proposing a novel framework. The proposed framework stores data in a hybrid database, supports multiple meta-models with model adapters, and uses among other methods truth discovery for conflict resolution and prioritization of data from certain sources.

The applicability of the proposed framework is shown in a study generating three different models using ArchiMate (enterprise modeling), securi-CAD (holistic threat modeling), and FTA (Fault Tree Analysis modeling). The study uses empirical data gathered from a virtualized Supervisory Control and Data Acquisition (SCADA) environment that depicts the systems and software setup of a small electric utility with a substation.

The remainder of the paper is structured as follows. Section 2 gives an overview of the related literature. Section 3 presents the proposed framework. Section 4 describes a study, where three different system models are automatically generated using the framework. Section 5 discusses the findings and future work, and Section 6 concludes our paper.

2 Related work

The following sections describe the related work on automated modeling and data integration with the focus on conflict resolution.

2.1 Automated and multi-source modeling

Enterprise systems architectures are complex. An organization uses various architectural elements such as information systems, business processes and technical components that have different relationships with each other [27]. Complex architecture also means that the management of the architecture becomes complicated. Enterprise architecture management needs to consider the impact that changes to an element have to the architecture as a whole. Enterprise architecture metamodels are often used to run different simulations and analyses [13,24], which requires granular and up-to-date information [33]. Up to date information is critical for the success of enterprise transformation projects [36].

The stakeholders involved in enterprise architecture management tend to be experts in their own fields and therefore need different viewpoints [42] of the enterprise to make decisions [25]. According to Niemi and Pekkola [35], artifacts – such as enterprise architecture models, need to cover multiple enterprise areas. Thus, enterprise architecture management is a complex process that needs granular and up to date information from various domains. Support from automation is therefore welcome to facilitate accurate and timely decisions.

Some authors have studied the requirements for automation in enterprise architecture management. Farwick et al. [9] investigated the requirements for architecture model maintenance through surveying. The authors found that the respondents were not happy with the amount of enterprise architecture items covered in the enterprise architecture repository. For collecting more items in time, automation was deemed important. The respondents prioritized correct data granularity, consistency, and being able to identify the source of data. Nowakowski et al. [36] found in their enterprise architecture management requirements study that optimal tools need to support heterogeneous data because enterprise data are available in different forms and spread over multiple repositories.

Various types of data are available for automation. Farwick et al. [10] looked into some potential information sources, measuring the data quality of the sources through actuality, completeness, correctness, and granularity. The most vital data sources for automatic data collection in an enterprise are found to be existing databases (including configuration management database), sensors in business process engines and application servers, and project portfolio

management tools. The main participants in the process of model creation are data providers, domain experts and enterprise architects [31]. Given the studied requirements and the multitude of available data sources, it would be preferable if an automated modeling solution uses more than one data source.

There have been attempts to generate models automatically from available data. Farwick et al. [11] propose a semi automated way of maintaining enterprise architecture models using technical interfaces and humans. The technical interfaces are assumed to be web services and the data used should be in a machine readable format. Alegria and Vasconcelos [2] propose a framework for IT architecture verification using passive network traffic analysis. The framework uses logical inference rules to infer facts about the architecture. The framework is implemented as a prototype and applied to a case study. The framework uses only a single data source and adding additional ones is listed as future work. Buschle et al. [5] generate an IT architecture model with the modeling language CySeMoL. The model is created by mapping Nexpose² vulnerability scanning results to the CySeMoL metamodel. Holm et al. [21] use a similar data set to populate an ArchiMate based enterprise architecture model. Both approaches use a single data source and leave some metamodel entities uninstantiated.

Veneberg et al. [48] propose a method to match operational data with enterprise architecture objects to support better decision making. The proposed method is suitable for creating enriched enterprise architecture models with six phases to solve a specific business problem. Enterprise architecture and operational data are required as input to the proposed solution. However, the solution does not describe how to do model-to-model transformations, which is important for model reuse and has been deemed an important aspect of model based system engineering by some authors [50]. Veneberg et al. regard model-to-model transformations future work.

Hinkelmann et al. [20] propose a new paradigm for the alignment of IT and business. In their approach, the alignment is done continuously in an agile manner. Hinkelmann et al. write that engineering an agile enterprise requires continuous design and redesign. Their approach is a version of the Plan-Do-Check-Act cycle. Enterprise models are used for identifying the adaption needs and for supporting change implementations. The authors combine machine interpretable ontology models with human interpretable graphical enterprise

² <https://www.rapid7.com/products/nexpose>

Table 1 Kühn's four integration patterns [26].

| Pattern | Description |
|------------------------|---|
| Reference pattern | Links exactly one element in one metamodel to another element in another metamodel |
| Extension pattern | Describes how one metamodel can be extended with concepts of another metamodel |
| Transformation pattern | Rules are used to create parts of one metamodel from another |
| Merge pattern | Specialization of the transformation pattern where two or more metamodels are used as sources |

architecture models. They use Kühn's patterns [26] for integrating different types of metamodels. We use these patterns to create transformation mappings between different model languages for the data source and the model adapters. The patterns are described in Table 1 and are part of the model adapter functionality that is described in Section 3.5.2.

To summarize, existing attempts to automate modeling have not attempted automated and generic multi-source multi-model creation. Moreover, some attempts have not covered model reuse at all and others have left parts of the metamodel uninstantiated. Many have used only one data source. In the proposed framework the authors of this paper address these issues.

2.2 Conflict resolution and evaluation of data sources

If data are collected from multiple heterogeneous data sources, information about the same object may differ between the data sources. This leads to conflicts that need to be resolved when the data are being merged into a single set. There are various strands of literature available on the topic, ranging from database schema matching to truth discovery.

Rahm and Bernstein [38] studied automatic schema matching, which is a basic problem in data integration of heterogeneous data. They claim that the structure and terminology of two schemas needs to be reconciled in order to integrate into one. Rahm and Bernstein propose a taxonomy that explains common schema matching techniques. The authors write that a match is possible for individual elements but also for combinations of elements. They also mention that most schema matching algorithms rely on auxiliary information, like dictionaries, previous decisions and user input. The authors stay on

a theoretical level and write that hopefully their ideas are useful for future implementations.

Haas et al. [16] study how to integrate diverse data from multiple sources where the data from each source is represented differently. The approach centers around database middleware systems. They conclude that for integration the data structure needs to be changed and the data must be cleaned so that duplicates and errors are eliminated. The authors combine schema and data transformation as a uniform mechanism to achieve data integration from multiple sources. Rundernsteiner et al. [41] survey how to integrate data from changing information sources. They find that data warehouses offer flexibility in a variety of environments, which could be monolithic, distributed or open ones. Data warehouses are built to handle autonomy, heterogeneity, and dynamicity of data sources. Tools such as filtering and merging of data from individual sources are used through a middle layer. In this type of integration approach, schema and interface changes are common. Our approach follows the basic principles described here, such as data cleaning and schema transformation, and the data are stored in a central database, although not a relational one.

Halevy et al. [18] survey bodies of work on data integration in large enterprises. They describe a multitude of different approaches such as automated schema mapping using linguistic similarities and overlaps in data types, schema mapping with machine learning based on a mediated schema, reconciling data on the instance level by automatically detecting references to the same object, adaptive query processing over unified architectures, model management that is using algebra for manipulating mappings and schemas, and semi-automatically generating semantic mappings for data integration systems. The authors conclude that data integration is a necessity and needs to be part of an organization's infrastructure. The work provides references to relevant methods, some of which have been tested e.g. model management for the proposed automated modeling framework.

Data provenance techniques are seen as a way to reduce ambiguity in data integration and enable data reuse by multiple authors. Simmhan et al. [45] propose a taxonomy of data provenance characteristics and demonstrate its use on scientific workflow approaches. Glavic et al. [15] present a categorization scheme for provenance models by surveying the existing approaches. They write that provenance is related to data annotation and temporal data management and versioning. Cheney et al. [6] describe three main database

provenance notations and compare their applications in confidence computation, view maintenance and annotation propagation. Our framework builds on simple data annotation and temporal versioning to reduce ambiguity in data integration.

Dayal et al. [7] look at how data from heterogeneous and distributed data sources is used for business intelligence. The work centers around a data integration pipeline, which needs a lot of manual effort when implemented using extract-transform-load tools. The authors find that traditional extract-transform-load tools do not provide support for collecting business requirements and fail to support optimized designs that meet quality and correctness requirements. The authors propose a layered methodology for the data integration flow life cycle to facilitate the design and implementation of optimal flows. The work is theoretic and addressed to the community at large.

While data warehouse based approaches focus mainly on data transformation and merging, there is also research directed at conflict resolution that considers the characteristics of data sources in a probabilistic and dynamic way. Some of these approaches look at the trustworthiness of the data sources and credibility of the data in conflict resolution. Truth discovery has emerged as a field to integrate multi-source noisy data. Li et al. [29] propose a framework to resolve conflicts in heterogeneous data from multiple sources using optimization. They claim that common approaches that are based on voting, where the data with the highest number of occurrences is seen as correct, are not good enough. These approaches do not take source reliability into account – if a data source is producing low quality information, it will affect the end result negatively. Instead, Li et al. propose an approach where a data source's reliability is estimated – a truth discovery model. The model is tested on weather, flight, and stock data. Adapting the model for more complicated conflict resolution scenarios is described as future work.

Li et al. [28] present a literature review on truth discovery models. They describe the main principles of truth discovery as the ability to estimate source reliability and the ability to select credible claims based on the reliability of each data source included in the comparison. For example, if a claim is supported by sources with high reliability then the claim will be assigned a high credibility value. The authors write that truth discovery algorithms make different assumptions about input data, relationships between sources and objects. Several of the truth discovery algorithms require data preprocessing before they can be used. Some of the differences are for example what is

allowed in the input data like duplicates or objects without conflicts. Most of the algorithms also assume that the input data are given in a way that is comparable – standardized format and textual representation.

Merging heterogeneous data is a challenge. Data warehouse principles help to clean, filter, and merge data, and data provenance principles help to reduce ambiguity about where the data originate from. Truth discovery methods offer a way to keep track of the data credibility in a probabilistic way and help to prioritize certain data sources in data conflict resolution [3]. To the authors' knowledge, no one has attempted to use truth discovery methods for automated modeling as we are proposing in our framework.

3 A framework for automatic model creation

This section describes our proposed framework for automatic model creation. First, a general overview is provided and then details are discussed in each subsection.

The purpose of the framework is to support automatic model creation of IT architectures using heterogeneous non-streaming data. The framework provides a structured way of addressing conflicts and filling gaps in the modeling data captured in multiple data sources. The framework relies on a common language and a common data structure. We define the common language as the language that describes the elements that are stored and the common structure as the way the elements are stored. The main components of the framework and the process of setting up the framework are shown in Fig 2.

The main process of automation follows a series of steps, which are denoted with corresponding numbers in Fig. 2. While we see this study as an independent contribution, it is a continuation of an earlier work [47].

1. The proposed approach is designed to be tool, metamodel, and model independent. The first step of setting up the framework is to decide what metamodels to use. More metamodels with similar semantics can later be iteratively added. Each metamodel used with the framework needs to be represented by a specification, which describes elements, associations between the elements, and properties. Section 3.1 discusses some potential metamodels.
2. The second step is to identify relevant data sources for modeling. These sources should cover the architecture layers that are to be modeled and

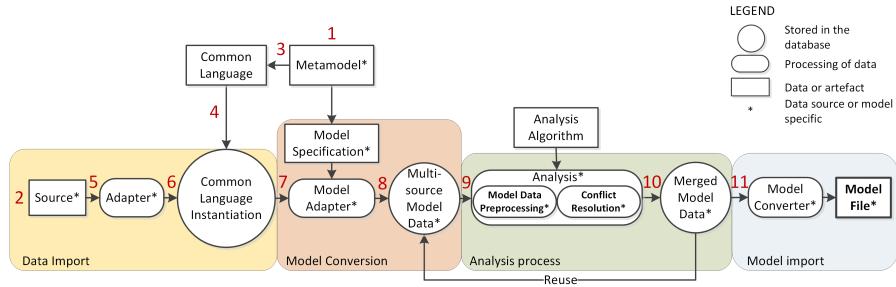


Fig. 2 Overview of the general process that consists of eleven steps.

could e.g. include configuration data, network scanners, and inventory systems. Section 3.2 discusses some possible data sources.

3. The third step is to create a common language and setting up a common data structure for storing the collected data. The common language needs to represent the architecture layers of interest. This involves choosing domain elements and the terminology for describing these. The metamodel specifications that the user wants to employ will be of help in creating the language and the structure. An example of a common language and a data structure is given in Section 3.3.
4. The fourth step is to set up a central database where the collected data are stored using the common language and data structure. Section 3.4 explains how a hybrid database can be of help in fulfilling this task.
5. The fifth step is to create adapters to transform the data from the language and structure of external sources to the common language and common data structure. Part of this process is identifying the data elements of interest in the imported files and mapping these to the common language elements. Adapters and data file formats are discussed in Section 3.5.
6. The sixth step is to store the transformed data in a central database and repeat the process regularly or whenever changes occur in the IT architecture, keeping temporal versions.
7. The seventh step is to create an adapter to transform the common language and data structure to the model of interest following the model's language and structure. The framework is designed to support multiple metamodels simultaneously. However, each metamodel needs a customized analysis to be set up that relies on common principles, but also takes differences in association multiplicities of metamodel elements into account. Section 3.5.2 describes how to set up a model adapter.

8. The transformed model is saved in the database separately in step eight.
9. The ninth step is to run the analysis in order to resolve conflicts in the data. Each data source is making a set of claims about the architecture and these claims might confirm each other or be contradictory. The analysis process includes the preparation of the data (preprocessing) by transforming the data into individual claims about the instances of metamodel elements, associations between the metamodel elements, and their properties. Once the claims are created, it is possible to reason over the credibility of these claims and the trustworthiness of the data sources making the claims using various algorithms. This paper focuses on solving the conflicts between the claims with a truth discovery algorithm. Section 3.6 describes the general analysis process.
10. The result of the analysis is a set of non-conflicting and unique claims about the architecture as the model structure. These claims are written back to the database in step ten, storing them for future analysis, consistency check and troubleshooting.
11. The final step is to create a file out of the analyzed data that can be imported into the modeling tool(s) of interest. This means using the import functionality of the chosen tool. Some common import formats are discussed in Section 3.7.

3.1 Supported models

The term metamodel is used for describing a model of a model, i.e. the model's structure and language. For the proposed framework, the metamodel(s) must meet a set of requirements. It must allow defining a metalayer and include the metaelements common in object-oriented languages. The structure needs to be based on classes (objects are instantiations of classes), associations between the classes, and properties of the classes.

The framework described here is meant only for data processing and visualization of data using a database system. Thus, a separate software tool for modeling is needed. Several metamodels and software tools are available on the market for IT architecture modeling. Three accessible and useful tools, which are used in the case study in Section 4, are (i) Archi for enterprise modeling using ArchiMate, (ii) securiCAD by foreseeti for holistic threat modeling and

attack simulation, and (iii) the Eclipse Modeling Framework based Fault-Tree Analysis (EMFTA)³.

3.2 Data sources

Different types of data are available in a typical enterprise system environment. Farwick et al. [10] and Johnson et al. [23] mention network monitors and scanners, configuration management databases, project portfolio management tools, enterprise service buses, change management tools, license management tools, directory services, business process engines, and release management tools as potential data sources. The trustworthiness of the data depends on the source type and collection frequency. For example, *configuration data* obtained directly from any system reflects a likely configuration of that system at that point in time, while data obtained using *network scanners* are somewhat less credible due to the dependence on a less accurate method called fingerprinting. Fingerprinting is a technique where UDP and TCP packets are sent to a host and the responses are compared to a database of known responses to determine the name of the operating system⁴ running on the host, or a service⁵. Since there are constant changes in the architecture the data need to be up-to-date in order for it to be useful [1]. Some of the data sources are able to provide constant data streams, e.g. syslog solutions. This framework does not support the collection of streaming data at the moment.

A non-comprehensive list of data sources is presented in Table 2.

3.3 Common language and data structure

The framework relies on (i) a common data structure describing the way the data are stored and (ii) on a common language describing the elements of the architecture layers of interest. Both are needed to support automatic model creation of multiple models and need to be general enough so that transformation from one IT architecture metamodel to another is possible. Fig. 3 shows a common data structure as a meta-metamodel, meaning that the elements defined in the meta-metamodel are the building blocks of the metamodel.

³ <https://github.com/cmu-sei/emfta>

⁴ <https://nmap.org/book/man-os-detection.html>

⁵ <https://nmap.org/book/man-version-detection.html>

Table 2 Examples of possible data sources for system modeling.

| Type | Examples | Type of data |
|--------------------------------|---|--|
| Network scanners | Vulnerability scanners, network topology scanners, packet analyzers | Hardware devices, software, vulnerabilities, network communication |
| Enterprise architecture models | Business, ICT architecture | Models of organization and its ICT architecture (in different views) |
| System management | Change, release, license management, directory services | System, network, process state |
| Events | SNMP, syslog implementations | System, network events |
| Security monitors | IDS, IPS, firewalls, SIEM solutions | System, network, process security |

The common data structure is currently restricted to metamodels following a general conceptualization of an object-oriented language. The main components of the common data structure are classes, associations, and properties. Each of those metaelements has a set of properties of its own, which are designed to support the definition of a common language and provide flexibility in data analysis. Classes represent IT architecture components, which are defined in a common language.

The IT architecture classes have four attributes – class (name), kind, label, and existence. The class attribute indicates the main class of the object and kind designates the subtype. For example, a class name could be *Host* and its subtype *Embedded*. If the exact kind of the architecture element is not known, the kind attribute can be left undefined. The label attribute contains the value of the most important property, which helps to uniquely define the particular class of object. An example of a label for the class *Host* is the *NetworkAddress*. The additional property values for a system class are stored as separate property objects, each linked to the system object. These property objects have their own kind and label attributes.

Classes are connected to each other through associations. These associations define how the component classes in the model relate to each other. An association class has the attributes label, target, and multiplicity. The label attribute shows the name of the association, target shows the direction of the association, and multiplicity shows how many target connections are allowed. For example a *Host* could be connected to many *SoftwareInstances* while it has only one *OperatingSystem*.

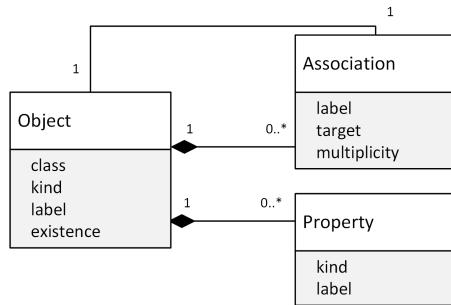


Fig. 3 Common Data Structure as a meta-metamodel.

The role of the common language in the framework is to represent the elements of interest and define the relationships between the elements. A common language is needed to standardize the input from different data sources and can be based on the most expressive metamodel of interest for simplicity. The data needed for other metamodels can then be derived from the more detailed one. It should include the dependencies between the architecture components and support using data sources that cover a subset of the common language elements.

The steps to devise a common language are:

- decide what metamodel is the most expressive for the domain of interest;
- capture all the key classes, properties, and associations of that metamodel;
- add any elements from other metamodels that were not expressed by the most expressive one.

The version of the framework described in this paper has been used to generate technology and partly application layer models (as opposed to e.g. business layer) and thus the common language of the framework is representative of these domains.

3.4 Back end database

The framework needs flexibility to store a variety of data types and large amounts of data. As the data requirements are evolving, the schema and the queries become more unpredictable. Storing this data in files is cumbersome as files lack the support for associating data elements and every functionality needs to be implemented manually. Traditional relational databases on the other hand are too restrictive. Relational data models need to be well defined

and implement integrity checks. Null values need to be handled consistently and doing this can be cumbersome. Therefore an alternative system for storing data is needed. Hybrid databases meet our flexibility and data scalability requirements.

The hybrid databases studied for the proposed framework support storing collections of JSON⁶ type nested documents of arbitrary and varying formats. Relationship modeling is supported with graph elements (edges) or links between documents. The structure of the data in the studied databases can be changed effortlessly. A database schema does not need to be defined although having a predefined structure simplifies constructing queries and data quality checks. Two database systems were found suitable for the framework due to functionality requirements and licensing terms, namely OrientDB⁷ and ArangoDB⁸. Both of these hybrid database systems have built-in support for graph visualization, which allows quick data quality control of the imported data.

To simplify queries and enable consistency checks, a database schema is implemented. It is implemented in a treelike manner. The data are stored in separate collections and metanodes have been added to support constructing graphs and queries. The metanodes also allow storing information about the data sources and the data in the database. These metanodes are *CollectedData*, *TimePoint*, and *Source*. The role of the *CollectedData* metanode is to tie all the graph elements together and therefore there is only one such metanode. It is the starting point of the graph. The creation time of each collected file is stored as a *TimePoint* metanode, which is linked to a *Source* metanode that designates the name of the particular data source (of the collected file). The *Source* metanode in turn is connected to the collected data (from the collected file), which are stored following the common data structure and common language notation. There can be multiple instantiations of the common language from different sources at different time points. Usually each data source covers only a part of the common language, thus multiple data sources of different types are needed to reach completeness in the sense of IT architecture metamodel coverage.

The data are written back to the database at three separate occasions during the process shown in Fig. 2. In the first two cases no merging is done

⁶ <http://www.json.org>

⁷ <http://orientdb.com>

⁸ <https://www.arangodb.com>

and each data source (at a time point) has its own subset of data even after model-to-model transformation. The data are merged only after the analysis step (Step 10 in Fig. 2) and then saved back as one merged set. The data at all three steps is kept for quality assurance and troubleshooting reasons. Data consistency is ensured by consistency checks. Each claim can be tracked by a unique identifier. The implemented consistency checks count conflicting and non-conflicting claims after different operations, check for duplicates, missing and irregular values. An irregular value might be for example a data flow without any client or server software association.

3.5 Adapters

Two different kinds of adapters are needed as described in Section 3. The first type of adapter is for transforming data from the format of the data sources to the common data structure and language (step 5 in Fig. 2). The second type of adapter is for transforming the common language format to a specific model language for a model language based analysis (step 7 in Fig. 2).

3.5.1 Data Source Adapter

Holistic systems and software data are available in different file formats from various sources. Some data can be exported as a file, which is the main input type for the proposed framework. The role of a data source adapter is to transform data from the source format and language to the common language and data structure. For that purpose a mapping between the data source file structure and the common language is needed. An example of such a mapping is described in the case study section of this paper and has also been covered in previous work [5, 21].

After the mapping between the input format and the common language has been created, the captured data can be parsed and stored in the back end database for further processing. Depending on the data format of the file, data format specific parsers like XML or regular expression based parsers can be used for extracting information. For example HP⁹ and Cisco¹⁰ devices allow to save running configuration files in text formats which can then be parsed with regular expression matching. The data gathered with network scanners

⁹ <ftp://ftp.hp.com/pub/networking/software/MgmtOct2005-59906023Chap06.pdf>

¹⁰ http://www.cisco.com/c/en/us/td/docs/ios/12_2/configfun/configuration/guide/ffun_c/fcf007.html

like Nmap can be exported using XML¹¹ and JSON syntax. The stored data from each data source includes the metadata describing when it was collected and also source details.

An important functionality for the data source adapters is checking whether a file has already been imported to avoid duplicating data. One way to do this is to create a hash of each file and when importing files, to check for the existence of that hash in the database and import the file only when the hash does not exist. Once the data from a file has been stored in a database, metanodes *TimePoint* and *Source* can be used to determine what data source at what time point is included in the model creation process. The assumption of the framework is that data from each unique data source (like Nexpose) is used only at a single point of time, in order to not bias the conflict resolution.

3.5.2 Model Adapter

Model adapters are needed if more than one type of model is to be automatically created with the framework. The role of the model adapter is to map the data from the common language to the modeling language and to account for the differences between the common language and the metamodel (class names, allowed associations between classes, and multiplicities of those associations). Kühn et al. have devised a pattern system for metamodel integration and in their paper [26] they describe three different types of integration patterns: (i) loose integration patterns, (ii) intermediate integration patterns, and (iii) strong integration patterns. The patterns are summarized in Table 1. Their work can be used as inspiration to devise a model adapter.

An implementation of the model adapter relies on the information about the metamodel and the methods that transform the source data to the model format and language. The information about the metamodel needs to be defined as a model specification. The specification needs to contain information about three things:

1. Model classes. E.g. *Network* or *Router*.
2. Associations and their multiplicities between the classes. For example, a *Router* has a one-to-many relationship with a *Network*, i.e. there can be many routers in a network. An *OperatingSystem* on the other hand has a one-to-one relationship with a *Host* as a host can have only one running operating system at a time (if we don't consider virtualization).

¹¹ <https://www.w3schools.com/xml/default.asp>

3. Dependencies between the classes. E.g. to be able to reason over the existence of *SoftwareProducts* they need to be defined in terms of the *Host* they are running on, thus every claim about the existence of a software product should contain information about the host the software is running on. A chain of dependent classes could be: (*Network*)(*Host*)(*Software*). A model specification needs to capture all these dependencies between classes.

The adapter changes the class names, the type and amount of properties the classes have, and associations between the classes. The transformation, as it is implemented in the framework, uses a separate function to transform each model class. Depending on the metamodel, no derivation is sometimes needed as there is a one-to-one match between the language and metamodel elements. At other times multiple common language elements might be combined into a single one, for example a network interface and a host in a common language might be represented in the metamodel only as a *Host* with multiple links to different *NetworkZone* objects.

The model adapter is implemented in the following way.

1. First a mapping between the common language and the model elements is created manually. The model specification is needed as input to the mapping process.
2. Secondly, a general function is created to query common language elements of a certain type from the database. Another general function is created that gets the associated elements of a certain type for an individual common language element using the element's identity value.
3. Thirdly, a separate transformation function that uses the general functions is created for each model class. The task for the transformation function is to transform every common language element to the appropriate metamodel element(s). Each transformation function is designed using the mapping logic between the common language elements and metamodel elements. The transformation of related associations and properties is included in the transformation process.
4. The fourth function joins all the individual model classes from the transformation functions and the fifth function saves the transformed and joined data back to the database as a separate collection.

During the transformation, the data needed for a particular model class is queried from the back end database in the chain of connected graph elements, class properties, and associations. To keep the uniqueness of the collected data

elements and allow traceability, unique identification numbers that identify each graph element are kept and reused. The associations between the elements are retained using the same unique id numbers.

3.6 Analysis

The core module of the framework is the analysis that is described in this section. The purpose of this analysis is to resolve conflicts in the data collected from multiple sources and to merge these into one coherent data set. Conflicts arise from the allowed multiplicities of associations between different model classes as defined in each model specification. For example, if the allowed association between some model classes is one-to-one and more than one connected class value exists, there is a conflict that needs to be resolved. Associations and conflicts are further explained in 3.6.4.

In the following parts of the paper the terms claim, entity, and value are used. *Claim* means a statement that a source makes about some element, which can be about a model object, model property or a model association. A claim consists of an entity – the target that the claim is about – and the claim's value. An *entity* is a unique model element and the *value* is a data value taken from the collected data that characterizes a particular model object, such as a network address for a host. We call this chosen data value the default (primary) property value. Default properties and associations are included in the entity representations implicitly (explained in 3.6.1). If there is more than one property or association, then these too are represented as individual claims.

It is the data sources, as explained in Section 3.2, that make claims. These claims from the data sources are either non-conflicting or conflicting. Imagine for example that one source claims that an operating system Windows XP is running on a host with a certain network address (for example 192.168.1.10). Another source might claim that the particular host is running a different operating system – Windows 7. The target of the claim (entity) is in this case the operating system of 192.168.1.10 and the value of the claim the name of the operating system. If the model specification we use allows only one operating system name at a time, the claims are conflicting.

The analysis consists mainly of two parts: (i) creation of claims from the collected data and (ii) resolution of conflicts between the claims. The first part is called “Model Data Preprocessing” and the second “Conflict Resolution” as

showed in Fig. 2. The conflict resolution part of the framework relies on a truth discovery algorithm, and is modular. Other algorithms can be used and for the sake of comparison, a naive algorithm based on common sense is also presented. The framework supports one algorithm at a time.

The quality of data varies between sources and the reliability of each source can be seen as an unknown variable that needs to be calculated. Truth discovery algorithms infer the trustworthiness of a data source from the data it provides and the credibility of the data from the reliability of the data source. There are different truth discovery algorithms available, some of them are summarized by Berti-quille and Borge-Holthoefer in [3]. The chosen truth discovery model is Latent Credibility Analysis (LCA) [37]. LCA is a probabilistic model that considers that a source can be fallible and thus provide wrong information. The model assumes that the sources are independent. LCA relies on Expectation Maximization as the main technique to calculate the probability of a claim to be true. The goal is to find correctness expectation, maximizing latent variables using the reliability of the source making the claim. The version of the algorithm, which has been implemented in the framework, is presented as Algorithm 1. The algorithm compares the data from the sources and based on prior evidence concludes that, given disagreement between the sources, one source is right (trustworthiness = 1) and the other is constantly wrong (trustworthiness = 0). The algorithm averages the trustworthiness of the values that each source provides. These values are weighted by the certainty the source has in its assertions. The LCA algorithm can only reason over mutually exclusive claims.

The input data for Algorithm 1 are the set of sources making claims (S), set of data objects (strings) that the claims relate to (D), set of conflicting values (strings) for each of the data objects (V), confidence matrix of the values (W) with confidence that each data source has in its claim expressed as ($w_{s,d}$). There is also the parameter of the prior claim credibility values (β_1), which can be a uniform scalar or set separately for each claim to resemble the beliefs of the user. K denotes the number of iterations for calculation. The recommended number of iterations is according to [37] around 100. The prior trustworthiness (T_s) and credibility values (β_1) should remain between 0 and 1. The confidence values ($w_{s,d}$) should start from 0, but have no upper limit. If the data source has full confidence in a claim, then that value is 1 or bigger, if the data source says nothing about the entity, the value is 0. The parameters ($\beta_1, w_{s,d}$) influence the outcome of the calculations. They provide

Algorithm 1: SimpleLCA algorithm [3]

```

Data:  $S, D, V, W, \beta_1, K$ 
Result:  $S$  with  $T_s$ ,  $V_d$  with  $C_v$ 

1 Initialization:  $\forall s \in S : T_s \leftarrow .8;$ 
2 while  $K > 0$  do
3   repeat  $K -$ ;
4   foreach  $d \in D$  do
5      $C_{d_{sum}} \leftarrow 0;$ 
6     foreach  $v \in V_d$  do
7        $C_v \leftarrow \beta_1 \cdot \prod_{s \in S_v} (T_s)^{w_{s,d}} \cdot \prod_{s' \in S_v} ((1 - T_{s'}) / (|V_d| - 1))^{w_{s,d}};$ 
8        $C_{d_{sum}} \leftarrow C_{d_{sum}} + C_v;$ 
9       foreach  $v \in V_d$  do
10       $C_v \leftarrow C_{d_{sum}} / C_v;$ 
11    end
12  end
13 end
14 foreach  $s \in S_v$  do
15    $T_s \leftarrow \sum_{v \in V_s} C_v \cdot w_{s,d} / \sum_{d \in D} w_{s,d};$ 
16 end
17 end
18 foreach  $d \in D$  do
19    $trueValue(d) \leftarrow argmax(C_v);$ 
20 end

```

an opportunity to fine-tune the input data according to the prior trust the user has in the sources. The outputs from the algorithm calculations are credibility values (C_v) for each of the conflicting claim values (V_d) and trustworthiness scores (T_s) for each of the data sources (S) making the conflicting claims. The algorithm uses the functions *trueValue* and *argmax*. The role of the *trueValue* function is to create a set of non-conflicting claims (represented with d), where each claim has the highest possible credibility value (C_v). The function *argmax* is used to identify the claim with the highest credibility value from each set of conflicting claims.

The preprocessing of data and conflict resolution are based on the following steps, which are further explained in the coming subsections.

1. *Obtaining transformed model data.* A set with the claims from all the data sources is obtained from the back end database. The set contains likely conflicting claims.
2. *Generating new unique identity values.* Identical claims from different sources need the same identity values when the data are merged, therefore identical

claim values need to be grouped and a common identity value needs to be assigned to each group.

3. *Merging data into long strings.* Preparation of claims for truth analysis by turning the claims into long strings that represent claims with all the elements (classes) involved. Long strings needed for textual string comparison in the conflict resolution analysis.
4. *Dividing claim data based on association multiplicities.* The model specification is used to classify elements based on their association types. One-to-one association means that there is only one possible right value and one-to-many means that many (all) values can exist.
5. *Generating negative values.* Negative values are generated based on each source's scope, claimed classes, and user specified data source capability, in order for the probabilistic truth analysis to express the likelihood that a value might not exist. For other types of analyses (like naive) negative values are not needed.
6. *Separating conflicting from non-conflicting claims.* Part of the analysis is identifying conflicts in claims and solving them. Non-conflicting claims are used without any further processing.
7. *Running truth analysis on conflicting claims.* Our framework uses truth analysis for conflict resolution, but it also supports other approaches. To demonstrate this, a naive approach is also introduced.
8. *Storing the results.* The results are stored in the back end database so that they can be, at any point of time, transformed into a model file.

3.6.1 Obtaining transformed model data

Before analyzing data, data first needs to be queried from the database. Classes and their dependencies in the model specification are used to construct the queries. For example, a network zone might be defined only using the class *Network*, while an operating system installed on a particular system is defined by the sequence of connected objects *Network*, *Host*, *SoftwareProduct* and their values.

The associations between metamodel elements are represented in different ways, depending on whether there is one or there are several named associations between the elements. If there is only one, the association is first stored between the connecting metamodel elements (as shown below) and later stored implicitly in long string format. We call the associations that are stored implic-

```
('Network', 'Host', 'Service'):
[[{'class': 'Network',
  'existence': None,
  'id': 'collectedModel/15291713',
  'label': '192.168.109.0',
  'source': ['2017-04-25_Nmap', 'subnet']}],
 {'class': 'Association',
  'target': 'Host',
  'kind': 'Connection'},
 {'class': 'Host',
  'existence': None,
  'id': 'collectedModel/15293715',
  'label': '192.168.109.21',
  'source': ['2017-04-25_Nmap', 'subnet']}],
 {'class': 'Association',
  'target': 'Service',
  'kind': 'Application execution'},
 {'class': 'Service',
  'existence': None,
  'id': 'collectedModel/15294101',
  'label': '49157-tcp-ncacn_http',
  'source': ['2017-04-25_Nmap', 'subnet']]], ...]
```

Listing 1 A service object with the associated objects as queried from the database.

itly *default* associations. Default associations are chosen based on if they define a metamodel element in a unique way in connection to the associated elements. If there is more than one association and many possible options for the default association, one association needs to be chosen and defined in the model specification. The rest of the associations we call *non-default*. Each non-default association is stored completely separately in a general group that contains all the additional associations. The definitions of additional associations contain the source and target information, and are linked to model objects through unique identity values (id values). This unique identifier represents a group of identical claims originating from different data sources. How this unique identifier is generated is explained in the next section.

Listing 1 shows an example of a service object from the group of service objects. The service object is defined through a sequence of dependent objects *Network*, *Host* and *Service*. For troubleshooting purposes, each dependent object contains an id and source information.

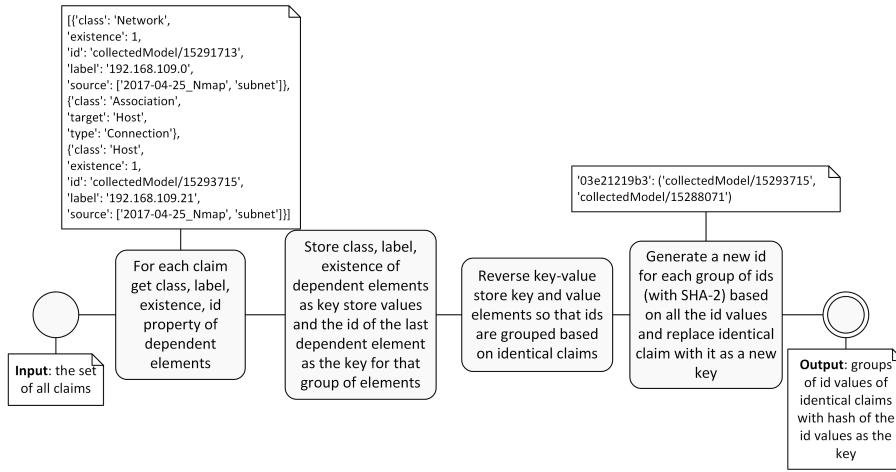


Fig. 4 Activity diagram of an algorithm for grouping id values of identical claims.

3.6.2 Generating new unique identity values

Multiple data sources are making conflicting and non-conflicting claims. Each of the data elements stored in the database is assigned a unique identity by the database. These unique identities are data source-specific. For the purposes of merging the data, new unique identities are needed that are not data source-specific, but are assigned based on the likeness of different claims. A method in the framework that consist of two steps, fulfills this task. An activity diagram of the algorithm that the method uses is shown in Fig. 4.

The first task of the method is to filter out the elements of the claims that will be compared for similarity (data source information is excluded) and then store them as a new data structure. This structure uses keys that define related values (see key-value stores). Each claim is stored with the database identity as the key. The second task of the method is to turn the data structure elements around – keys into values and claims into keys. This way the database identities (that are data source specific) are grouped based on similarity of claims (entity, value, existence). Here the entity means the main target of the claim like a service, the value means the value assigned to the entity like the name of the service, and existence is a user assigned (or inferred from the data) likelihood that the claim exists. Then the method replaces those claims with a new hash identity that is generated using database id values.

Listing 2 shows an example of a group of id values of matching claims from different sources. The key of the group is a truncated SHA-2 hash value that

```
'03e21219b3': ('collectedModel/15294101', 'collectedModel
/15288071'),
```

Listing 2 A group of id values of identical claims from different sources.

```
('Network', 'Host', 'Service'):
[{'class': 'Service',
'existence': None,
'fullName':
'<<Network>>192.168.109.0<<Host>>192.168.109.21<<Service
>>49157-tcp-ncacn_http',
'id': ['collectedModel/15291713',
'collectedModel/15293715',
'collectedModel/15294101'],
'label': '49157-tcp-ncacn_http',
'scope': 'subnet',
'source': '2017-04-25_Nmap'},
```

Listing 3 A transformed claim about a service running on a particular host.

has been generated using the id values, which are in the format collectedModel/... . The first part of the original id shows the database collection name and the second is the unique id that each data element has in the database. There is no hash collision handling; this is planned as future work.

3.6.3 Merging data into long strings

The truth discovery algorithm is able to reason over data claims that follow a certain format. The queried data need to be changed into comparable text strings to be suitable for the truth discovery algorithm. The truth discovery analysis needs strings that represent an entity about what claim is being made, the claim itself (value), and then the source's name that made the claim. The source, in our case, is the source of collected data at a certain point in time.

An example of a transformed data element is shown as Listing 3. The example contains a claim about a service running on a particular host. The string that is created with this particular method is called fullName and contains classes and primary associations (implicitly) and the primary property for each class.

3.6.4 Dividing claim data based on association multiplicities

Each data source makes a claim about the value of some model object (entity). The claims made can be mutually exclusive, meaning that at most one value in the set of claims about one entity can be true. The claims can be also mutually non-exclusive, which means that all values can be true or false independently of each other. For example, a mutually exclusive claim about a node is that it runs a certain operating system. So if one source says that the node runs Linux and another source says that it runs Windows, both cannot be correct at the same time (though dual boot allows the truth values to change over time). An example of a non-exclusive claim is that a certain software application is installed. So if one source says that the node runs Emacs and another that it runs Apache, there is no contradiction and both can be true or false independently.

The exclusiveness of the claim depends on the association multiplicity between two model classes. If the multiplicity in the metamodel is one-to-many, then this is sufficient for the claims to be non-exclusive, for there could be many different objects (class instantiations) associated with the object at hand. As in the example above, many pieces of software could run on a single node, and the multiplicity of the “runs-on” relation is thus one-to-many. However if the association in the metamodel is one-to-one, it is unfortunately *not* sufficient to know that differing claims from different data sources are mutually exclusive, at least not under a primitive string comparison notion of differing claims. For even though the string-different claims “Windows” and “Linux” are mutually exclusive, the string-different claims “Windows” and “Windows 7” are not. To solve this complication, a refined notion of differing claims is needed. We believe that this could be accomplished using ontologies, and aim to do so in future work. However, in this paper, we disregard this difficulty and proceed as if one-to-one multiplicities were sufficient to establish the mutual exclusiveness of claims. For the data we use, this is a reasonable simplification, as will be seen in Sections 4 and 5. However, there is no guarantee that this simplification will always be reasonable. It depends, among other things, on the particularities of the data sources used and the system analyzed.

In conclusion, the task of this step of the analysis is to divide the model data into two sets based on the multiplicities of associations defined in the model specification, where one set contains exclusive and the other non-exclusive claims.

3.6.5 Generating negative values

There are truth algorithms that only can reason over the credibility of mutually exclusive claims, like LCA. That means there needs to be one or more subsets of mutually exclusive claims, where only one claim in each subset is determined to be true. There are also algorithms that are not limited to mutually exclusive claims and can find more than one claim in a set of claims to be true. The Latent Truth Model (LTM) [3] is such an algorithm. The problem with the LTM type of algorithms is that they treat all claims as mutually non-exclusive. While it is not possible to restrict LTM to reason over the credibility of a mutually exclusive set in a way that only one claim is true, it is possible to use LCA to reason over mutually non-exclusive claims with a small modification – the addition of negative claims.

The authors of this paper preferred to use a single truth algorithm to reason over the credibility of the whole data set instead of using multiple algorithms simultaneously, because the truth discovery models behind the algorithms are based on different assumptions and their results might not be comparable. Therefore the simple version of the LCA is chosen as the main truth algorithm. To be able to use the LCA algorithm, each mutually non-exclusive claim needs to be treated as an individual claim that can have a positive or negative outcome. For that purpose additional negative claims need to be added to the data set. Negative claims can be added for each mutually non-exclusive claim that meets certain requirements, which are explained below. While the generation of negative claims is needed to transform a mutually non-exclusive data set into a mutually exclusive set, for consistency reasons, negative claims should also be added to the already mutually exclusive claims.

A negative claim expresses the likelihood that an object (alternatively value) might not exist. A negative claim must be made by a data source that has the potential to make it. To be able to generate negative claim values for a data source, certain criteria need to be met. The data source needs to have the capability to make a positive claim although it is not doing it. The capability is defined first by scope (assigned to each source), which can have two values – “system” or “subnet”. “System” means that the data source has knowledge only about one particular system. Many configuration files fall under this category. “Subnet” on the other hand means that the data source has knowledge about other nodes on a particular subnet. For example, network scanners. The second criterion is the data source’s capability to make claims

about the model class in question. If the source makes other claims about a particular class of elements, then it is assumed that it has that capability. The third criterion is the similarity of data sources. The data sources need to have similar capabilities to find information, to be able to qualify for negative claim making. For example Nessus and Nmap qualify for making negative claims about claims from Nexpose. The similarity of the data sources included in the analysis needs to be decided by the user and defined beforehand manually.

Tables 3 and 4 show how the entity values could be complemented in both cases. For example, that a router has a firewall installed (Case One) or a computer has different software programs running on it (Case Two). A negative value designates the non-existence/non-credibility of a particular entity. As explained, a negative value is generated only if there is a data source that meets the three criteria.

Case One describes a one-to-one relationship where only one claim about an entity can be true. However, as any claim means that the entity exists, the real question is if there is any source that meets the criteria to claim it does not exist. In this case ‘Data source 4’ meets the criteria and thus a negative value is generated. Case Two describes a one-to-many relationship. All of the claims can be either true or negative (designated as false), thus the claims need to be treated like the entities in the first case. Simply put, the claims about an entity are divided into smaller groups where each claim becomes an entity itself. If there are data sources that meet the criteria to be able to make negative claims, the negative claims are generated, otherwise the claim is considered to be true (non-conflicting) and moved to the results data set.

The algorithm for negative value generation function is shown in Fig. 5. The method using the algorithm checks for each claim in a set of claims if there are other sources that meet the three criteria (class, scope, source similarity) and are not making claims about the target of the claim (entity) as shown in Tables 3 and 4. If the criteria are met and no claims are being made, a negative claim with the new source information is generated.

3.6.6 Separating conflicting from non-conflicting claims

Once the negative claims have been generated according to the rules explained in the previous subsection and the format of the one-to-many claims has been changed, conflicting data (claims) need to be separated from non-conflicting. The criteria for identifying a conflicting claim is (i) that more than one data

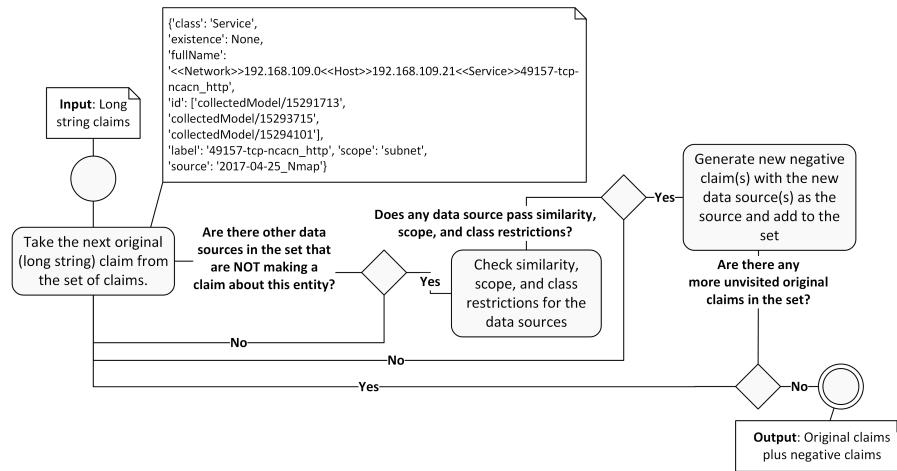


Fig. 5 Activity diagram of an algorithm for generating negative values.

Table 3 Case One: one-to-one relationship where only one claim about an entity can be true.

| Entity A | Data source | Value of entity A | Action |
|----------------------------|---------------|----------------------|-----------|
| <Router>10.1.1.1<Firewall> | Data source 1 | Stateful firewall | Exists |
| <Router>10.1.1.1<Firewall> | Data source 2 | Packet filter | Exists |
| <Router>10.1.1.1<Firewall> | Data source 3 | Application firewall | Exists |
| <Router>10.1.1.1<Firewall> | Data source 4 | Negative | Generated |

Table 4 Case Two: one-to-many relationship. All of the claims can be true or be negative (designated as false).

| Entities B, C, D | Data source | Value | Action |
|-------------------------------------|---------------|-------|------------------|
| <Host>WIN-012<Service>ArangoDB | Data source 1 | True | Already existing |
| <Host>WIN-012<Service>ArangoDB | Data source 2 | False | Generated |
| <Host>WIN-012<Service>RabbitMQ | Data source 3 | True | Already existing |
| <Host>WIN-012<Service>RabbitMQ | Data source 4 | False | Generated |
| <Host>WIN-012<Service>Support Agent | Data source 5 | True | Already existing |
| <Host>WIN-012<Service>Support Agent | Data source 6 | False | Generated |

source is making claims about an entity (object of the claim) and (ii) that these claims are different (have different values). The separation of claims into two data sets can be done by simply comparing the claims made by different data sources with each other (long string comparison). Ontologies could offer

a more powerful comparison than just string comparison, but this is not in the scope of this paper.

What gets compared is an entity and the values of an entity. An entity is for example a certain operating system running on a host and the value of the entity is the operating system's name. If there are more than one non-identical claim about the host operating system value, both claims are classified as conflicting. If there are no contradicting entity values for a claim, then it is classified as non-conflicting and stored. The stored non-conflicting claims are later merged with the claims from the truth analysis.

3.6.7 Running truth analysis on conflicting claims

Conflict resolution is possible once the data has been preprocessed, transformed into claims, and sorted. The chosen algorithm for conflict resolution – LCA, is presented as Algorithm 1. However, other approaches are also supported by the framework and for comparison, a naive algorithm is introduced in this subsection.

The naive algorithm is based on common sense. It does not support using prior probability values to describe the data sets and does not support using negative claims. Therefore only mutually exclusive claims can be included in the naive conflict resolution analysis and mutually non-exclusive ones are treated as non-conflicting data. The goal for the naive algorithm (see Fig. 6) is to find, for each entity (object) of the claim, the most frequently occurring value. If the most frequently occurring value cannot be determined because the conflicting values occur equally often, then two other steps are possible. Firstly, the claim from the most trustworthy source can be picked. The trustworthiness of a source is in this case determined by dividing the total amount of matching values by the total amount of conflicting values. These amounts are counted from the whole data set, not only from a particular claim subset. If the scores for all the data sources making the particular claim are the same, then instead a random value is picked.

3.6.8 Storing the results

The last step in the analysis process is storing the results in the back end database. Before the results are written to the database, the truth discovery results are merged with the non-conflicting data set. This way the chain of class dependencies can be followed while storing the data. The storing of claims is

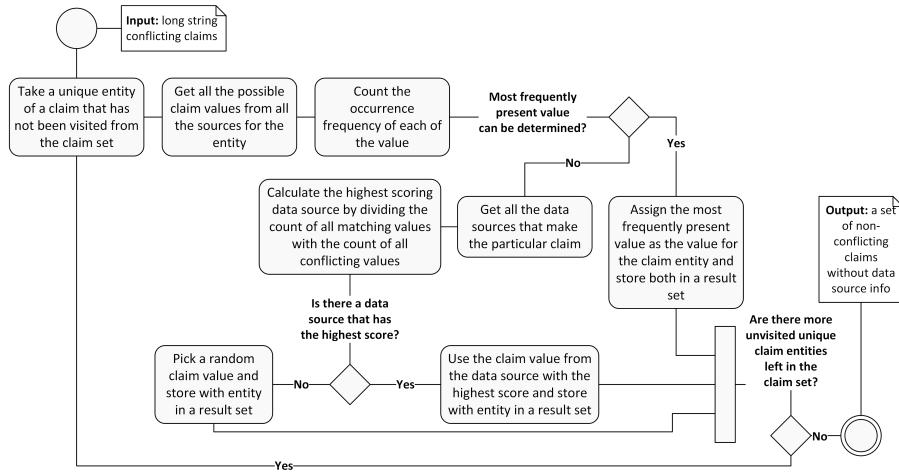


Fig. 6 Activity diagram of naive conflict resolution.

done starting with the objects that have the smallest number of dependent classes and gradually increasing the number. For example, all network zones are stored first, while the hosts connected to these zones are stored in the next iteration. This means that the claims that depend on the existence of some other claim can be excluded if the other claim is found to be negative. So, no host information would be stored if the network zones they depend on are missing. The results are stored in the database using the unique id numbers that are generated at the beginning of the analysis process (see Section 3.6.2). This way the correct associations between model elements are guaranteed.

After the merged data have been stored in the database using the chain of dependence, the data are now ready to be imported to the corresponding modeling tool. The stored results can also be reused in future analysis processes as the ground truth.

3.7 Model-to-model transformation

The modeling tools considered either support import with the tool's proprietary format, or the modeling tool files containing the models can be generated directly. The prevalent file formats are XML syntax based. To create the import files or model files directly one needs to have the appropriate tool documentation available to do so.

4 Study: testing the framework on a SCADA environment

This section presents our implementation of the proposed framework and employs it in a study of a SCADA environment. The prototype of the framework is available for download¹².

4.1 SCADA environment

The purpose of Supervisory Control and Data Acquisition (SCADA) systems is to monitor and control large scale processes like industrial production, electricity distribution, or building automation systems. SCADA systems offer a link between (i) the embedded software in various sensors and actuators, and (ii) a remote operator control room. The typical setup consists of a master station, Remote Terminal Units (RTUs), and a communication system [4]. A master station (SCADA server) is responsible for communicating with the geographically dispersed field equipment. Data from the field is acquired through the system part called Front End (FE). Human users operate the SCADA system through a human machine interface (HMI) with different so called consoles. Operator consoles are used by system operators for controlling the supervised process. Engineering consoles (alternative name: maintenance consoles) are used for control system maintenance; database and situational picture maintenance. RTUs are the sources of measurements and states for a SCADA system.

The experimental setup was designed as part of the European Union FP7 project SEGRID¹³ and aims to represent the SCADA setup of a small electric utility. The central part of the architecture is a SCADA system running on four servers. Two of the servers are the main SCADA servers in a redundant setup that run on Red Hat Enterprise Linux. The control system includes functions for remote real time data acquisition and remote control of process devices. The servers also run electric grid specific tools for monitoring, control, and performance optimization of the grid and historians to store historic process data. There is a separate Windows 2012 Server for Active Directory and Kerberos¹⁴ based authentication. The fourth machine, also Windows 2012 Server, has data engineering software installed and an HMI. This machine is

¹² <https://www.kth.se/profile/robertl/page/automatic-data-collection-and-modeling>

¹³ <https://segrid.eu>

¹⁴ <https://web.mit.edu/kerberos>

also set to perform the tasks of a FE. With the help of the FE software the SCADA system regularly fetches measurements and sends control signals to one destination which simulates a substation through a wide area network. The substation is represented in the environment by an RTU.

The servers and the local network are running in a virtual environment, but the wide area network and the RTU are running on physical hardware. The wide area network has been set up using industrial wireless mesh routers. The network is separated using virtual local area networks (VLAN)¹⁵. An open source pfSense firewall¹⁶ is routing and blocking the traffic between different VLANs. The setup is shown in Fig. 7 and summarized in Tables 5 and 6.

In Fig. 7 DMZ stands for DeMilitarized Zone, EMS for Tropos Control Element Management System, ERP for Enterprise Resource Planning system, RTU for Remote Terminal Unit, and ICC for Inter-Control Center Communications.

The environment is representative of a small power utility and the software and equipment run are designed for this purpose. The network environment has been segmented to different functionality oriented zones to achieve conformity to real environments. A major difference from a real environment is that the systems are not connected to any real power processes, i.e. they have not been set up to control any cyber-physical system. However, the goal of the paper is to study data collection from systems and software (IT) environments, thus having a working power system is not required for meeting this particular goal. Another difference is that virtualization is used extensively here, however this does not drastically influence the data collection either.

4.2 Data collection in SCADA setup

Three different types of data are collected for model creation: (i) configuration data obtained from systems directly, (ii) data from network scanners, and (iii) network traffic. The following subsections provide an overview of the data collected and the methods used.

¹⁵ <http://www.ieee802.org/1/pages/802.1Q-2014.html>

¹⁶ <https://www.pfsense.org/download>

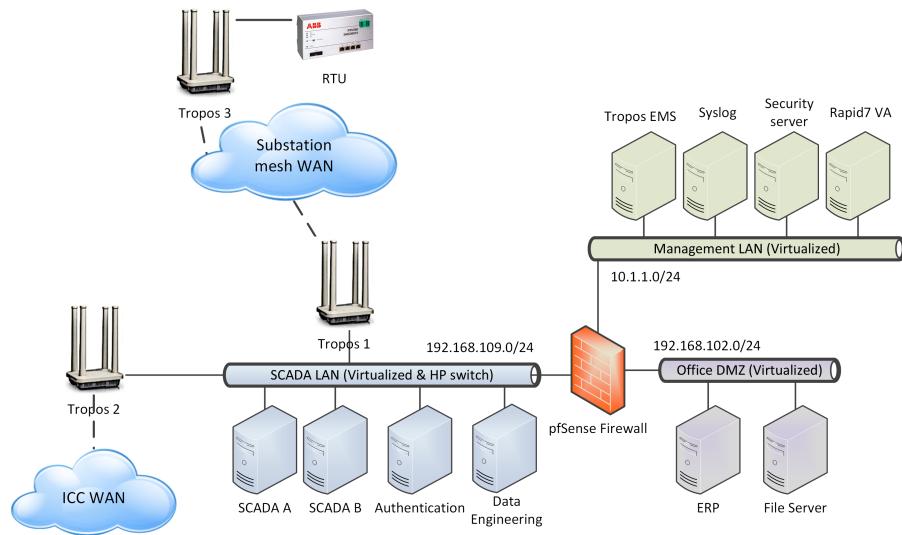


Fig. 7 A small utility SCADA environment setup.

Table 5 Devices in the SCADA environment with short descriptions.

| Network zone | Device | Description |
|-------------------|------------------------------|--|
| SCADA subnet | Authentication server (AD) | Domain controller and DHCP server |
| | Data engineering server (DE) | HMI, Frontend system |
| | RTU | Endpoint to substation |
| | SCADA server A | SCADA, NMS application, historian |
| | SCADA server B | SCADA, NMS application, historian |
| Office subnet | ERP server | Office application for resource management |
| | File server | File server for data from SCADA zone |
| Management subnet | Syslog | Syslog server to consolidate logs |
| | Rapid7VA | Vulnerability scanner Nmap |
| | Security server | Network scanners and passive traffic capturing |
| | Tropos EMS | Central management for Tropos routers |

4.2.1 Configuration data

The configuration data in the lab are collected from general types of operating systems and network devices. The data set obtained from Windows and Linux based operating systems include information about installed software programs, running services, open ports, and configured users. The configura-

Table 6 Network devices in the SCADA environment.

| Device | Network zones | Description |
|---------------------|-------------------------------|---|
| pfSense | Management, office, SCADA | Router and firewall between different subnets |
| Tropos Wifi routers | Management, SCADA, substation | Substation wide area networks and subnets |
| Vmware switches | Management, office, SCADA | SCADA, Office and management subnets |
| HP Procurve switch | SCADA | SCADA subnet extension for compatibility reason |

Table 7 Collected configuration data and collection methods for Windows²⁰ and Linux²¹.

| Operating system – Windows 2012 Server | |
|--|---|
| Method | PowerShell cmdlets (PS) |
| Data | Installed Software, running services, local user accounts, open ports |
| Operating system – Red Hat Enterprise Linux 7 | |
| Methods | getent, netstat, systemctl, yum |
| Data | local users, open ports, services, installed software |
| Firewall – pfSense | |
| Method | Backup file |
| Data | interfaces, IP address ranges, running services, firewall rules |
| Switch – HP Procurve | |
| Method | Running configuration file |
| Data | VLAN ports and IPs |

tion data set collected from the HP Procurve switch and the pfSense firewall gives an overview of the network zones and management of IP addresses. Some of the devices in the lab use proprietary data formats for storing configuration data and do not support exporting this data into a machine readable format. This problem relates to for instance ABB Tropos¹⁷ routers. Table 7 gives a brief overview of the methods used and data obtained.

¹⁷ <http://new.abb.com/network-management/communication-networks/wireless-networks>

¹⁸ <https://technet.microsoft.com/en-us/library/dd772285.aspx>

¹⁹ <https://highon.coffee/blog/linux-commands-cheat-sheet>

²⁰ <https://technet.microsoft.com/en-us/library/dd772285.aspx>

²¹ <https://highon.coffee/blog/linux-commands-cheat-sheet>

4.2.2 Scanner data

Three different scanners are used for examining the SCADA environment. Two vulnerability scanners, Nessus²² and Nmap²³, and the network scanner Nmap²⁴. We run Nmap on a separate virtual machine, Rapid7VA, while Nmap and Nessus are installed on a Microsoft Windows server that has access to all the network zones. The scanners use a variety of methods for obtaining information about the system scanned, including pattern recognition. They are able to authenticate themselves to certain systems in order to obtain detailed configurations. The information collected includes installed software, the patch level of the software, open ports, existing vulnerabilities, and potential configuration problems.

4.2.3 Other data

Two network capturing tools are used to gather network traffic information. Wireshark²⁵ can capture network traffic, allows filtering based on types of protocols, and can follow sessions. p0f²⁶ is a tool that does passive network traffic monitoring to identify the operating systems on both endpoints of a TCP connection. The tool can detect network address translation and connection sharing.

4.3 Applying the framework

In the following subsections the proposed framework is applied to find answers to the two research questions posed in the introduction. We study whether the framework allows to automate modeling when multiple heterogeneous data sources are needed and if the automation can be done in a way where data source trustworthiness is tracked in conflict resolution. For this, three different tools for holistic systems and software (enterprise) modeling are employed.

²² <https://www.tenable.com/products/nessus-vulnerability-scanner>

²³ <https://www.rapid7.com/products/nmap>

²⁴ <https://nmap.org>

²⁵ <https://www.wireshark.org>

²⁶ <http://lcamtuf.coredump.cx/p0f3>

4.3.1 Tools and metamodels

For the study three different metamodels are used – ArchiMate²⁷, Eclipse Modeling Framework based Fault-Tree Analysis (EMFTA)²⁸ and securiLANG²⁹. A short introduction to each metamodel is provided below.

ArchiMate is a technical standard from The Open Group for the description, analysis, and visualization of three interconnected enterprise architecture domains – business, application and technology. The technology domain which is the main focus in this study (together with software and data), deals with an enterprise's hardware and communication infrastructure. There are commercial and open source tools available that implement ArchiMate, e.g., BiZZdesign Enterprise Studio and Archi. In this study we use Archi³⁰.

EMFTA is a simple fault tree analysis extension for Eclipse. The extension supports fault tree probability analysis and visualization, and it is open-source.

securiLANG is a metamodel for quantitative threat modeling and attack simulations, and is implemented in a modeling tool called securiCAD³¹. It has its foundation in more than 10 years of dedicated research conducted by a full team of researchers at KTH Royal Institute of Technology. The core concept in securiCAD is the use of attack graphs, which are rendered under the modeling surface. As a user, you model the software and systems architecture, then the tool simulates all possible attacks. The tool has built in statistical expertise on how hard it is to make a specific type of attack, to a certain type of system, given the security features of this system. The key measure is Time-To-Compromise (TTC).

EMFTA and securiCAD are tools for simulation and unlike Archi they need evidence from the real world to run the simulations. This means that these two tools have additional data needs, over and above what is collected in the following case studies. The main role of Archi is to document existing architectures and communicate ideas for new ones.

²⁷ <http://pubs.opengroup.org/architecture/archimate2-doc>

²⁸ <https://github.com/cmu-sei/emfta>

²⁹ <https://www.foreseeti.com/products>

³⁰ <https://www.archimatetool.com>

³¹ <https://www.foreseeti.com/products>

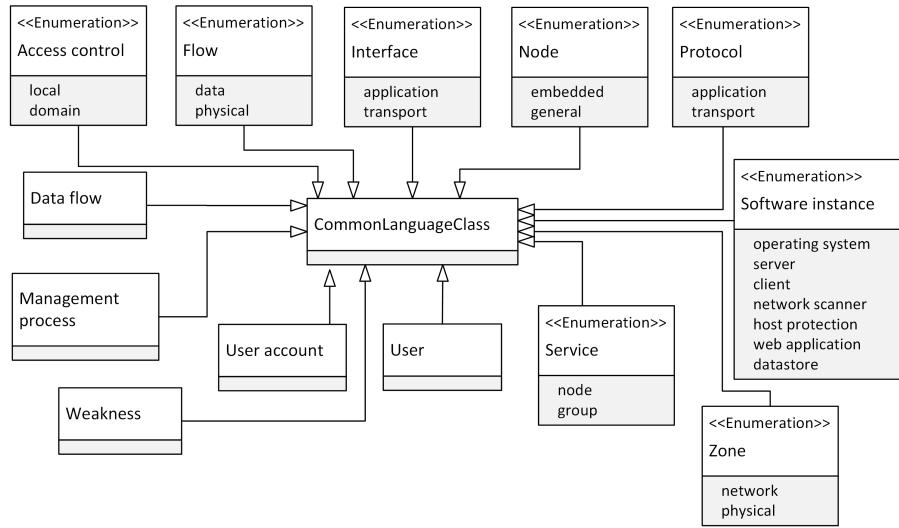


Fig. 8 Common language of the three chosen metamodels for the study of the SCADA environment.

4.3.2 Study: common language

The data collected need to be stored in a common language to resolve conflicts. This common language expresses the details of the models in consideration.

The three metamodels chosen for automatic model creation allow to analyze and document the systems architecture of and data flows within an organization. Thus, the common language we have developed for the implementation of the framework represents those elements. It describes IT architecture components like networks, hosts, and software, but also management processes, data stores and data flows. The language is developed by first identifying the most expressive metamodel of interest, which is securiLANG. After all the securiLANG elements are captured in the common language, the resulting language is compared with the ArchiMate and the EMFTA metamodels to see if it covers those metamodel elements as well, and is improved.

Fig. 8 shows some of the classes and kinds of the common language. The common language is implemented in a hybrid database and follows the common data structure described in Section 3.3.

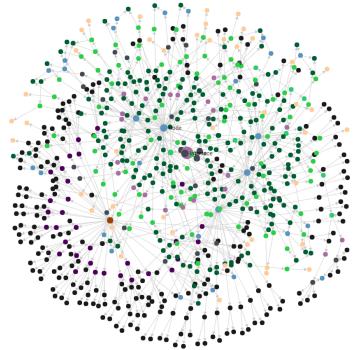


Fig. 9 Collected data as a graph. Each vertex represents one common language element. The colors show different common language classes.

4.3.3 Data Source Adapters

A total of eleven different adapters are used to import the data from the data sources in the lab to the database. The imported data are transformed to the common language. Some sources have data available at multiple time points, but only the most recent is used for each source.

4.3.4 Database

A hybrid database system called ArangoDB version 3.2 is used for storing the collected data. The choice was made based on functionality, licensing terms and software bugs encountered during implementation.

The data in the hybrid database are stored in collections as vertices and edges and can be traversed as a graph³². ArangoDB has a built in simple graph visualization interface, which can be used for quick quality control of the stored data and discovering patterns in the data. Fig. 9 shows the graph visualization of some of the data collected (figure size limited by the user interface). In the figure different element classes have different colors.

4.3.5 Analysis and Results

This section describes the creation of the three models that were introduced in Section 4.3.1. This activity involves model conversion, analysis and model import (see Fig. 2). The end goal is to get a separate IT architecture model for each of the three modeling tools, securiCAD, ArchiMate and EMFTA.

³² <https://docs.arangodb.com/3.2/Manual/Graphs>

The purpose of the model conversion part of the framework is to prepare the collected data set for model specific analysis. The model conversion starts with the transformation of the data set that is in common language to the three model languages in focus, securiLANG, ArchiMate and EMFTA. For that purpose model language specific model adapters need to be configured and run. The model adapters use model specifications to support the transformation process. A model specification contains the definitions of elements, associations, and dependencies between the elements. The central elements for example in securiLANG are networks and hosts. A *Host* object typically has connections to other objects such as *Network*, various types of software, and *AccessControl*. Unique *Hosts* are defined through the chain of dependency including the connected *Network* object while instances of software are defined through connections to *Host* objects.

The transformation logic of all three model adapters follows Kühn's integration patterns presented in Table 1. The adapters are implemented in a way that there is a separate function for the transformation of each of the model classes. For some common language elements there are direct equals, so the "Reference pattern" can be followed. In securiCAD's case *Network* objects belong to this category. For other model elements transformation rules ("Transformation pattern") are needed in order to transform multiple common language elements to a single metamodel element. For example, *Hosts* are created combining the common language elements *Interface* and *Node*. There are also elements that do not always have direct equals like the class *Software Product*. It would be possible to generate *Software Product* elements automatically, as they represent the static versions of software instances, but that would skew the truth analysis results. Such a generation step should therefore, if at all, be done after the conflict analysis has taken place. A detailed explanation of model adapters is given in Section 3.5.2.

Because of the limited expressiveness of the ArchiMate model, the ArchiMate specification needs to be customized – two different versions of the class *System Software* are needed in the model specification. This is because *System Software* can represent both operating systems and installed software. The need comes from real life. Due to the way the lab has been set up we need to restrict the number of operating systems that can be modeled on one host. There are no restrictions to the amount of software a host can run, however. Thus, a device can be connected to only one *System Software* element that expresses the operating system software running on it while it can be

connected to multiple *System Software* elements that express other installed software. Another issue is that some of the ArchiMate model elements can express logical groups of other elements. For example, a *Node* can express the logical grouping of a *Device* and *System Software*. Therefore *Node* elements should be generated when *Device*, *System Software* and associations between them are known. To avoid this element generation influencing truth discovery analysis, *Node* elements are generated after the truth discovery conflict analysis has taken place. The reasoning behind adding new elements is based on logical deduction. While there is a possibility of accuracy problems, based on the authors' experience of using the framework, no such problems were evident.

The EMFTA metamodel contains a limited amount of elements designed only for availability modeling and the metamodel lacks technology domain knowledge. Therefore, the technology layer elements from the common language are used to complement the EMFTA model. The model adapter transforms the data from the common language to the EMFTA model for analysis. The EMFTA adapter makes no changes to the common language.

The purpose of the analysis part of the framework is to solve conflicts in the model specific data sets. Thus, the analysis part of the framework needs to be configured to support each metamodel. The analysis needs to use a correct model specification together with the right multiplicities and class names, and the analyzed data needs to be stored in the designated database collections. For truth discovery analysis, the capability of the data sources to negate each other's claims on certain conditions (as discussed in Section 3.6.5) needs to be configured as well. The capability configuration answers the questions like the following: If network scanner Nexpose claims that a certain host in a certain network zone exists, could Nessus, having scanned the same network zone, but not having detected that particular host, reasonably claim that the host in fact does not exist?

Once the analysis process has been configured the framework can be run, which means first running the preprocessing to transform data into claims and separate conflicts. The preprocessing is analysis algorithm specific. The framework supports two different algorithms (truth discovery, naive) for conflict resolution as alternatives to each other. Both algorithms are described in Section 3.6. The truth discovery analysis can be fine-tuned with parameters, as explained in Section 3.6.7.

Table 8 Claims from all the data sources included in the analysis. N/A denotes Not Applicable.

| Name | securiCAD | | ArchiMate | | EMFTA | |
|---------------------------|-----------|-------|-----------|-------|-------|-------|
| | Naive | Truth | Naive | Truth | Naive | Truth |
| Initial claims | 1477 | 1477 | 1283 | 1283 | 1045 | 1045 |
| Generated negative | N/A | 1334 | N/A | 1302 | N/A | 1249 |
| Conflicting claims | 84 | 2201 | 82 | 2138 | 0 | 1983 |
| Non-conflicting claims | 1393 | 610 | 1201 | 447 | 1045 | 311 |
| Conflicts solved claims | 21 | 738 | 20 | 705 | 0 | 637 |
| Final non-negative claims | 1414 | 846 | 1221 | 671 | 1045 | 448 |

During the preprocessing of the data for the truth analysis uncertainty is factored in as negative claims. No negative claims are used for the naive algorithm. After preprocessing, conflict resolution can be conducted. Eventually, the resulting data set from either of the analysis methods is merged with the non-conflicting data that was separated during preprocessing and then stored in the database as a model specific merged data set. The merged data sets from the database can then be imported to the modeling tools.

Table 8 shows the total number of claims processed for each of the three models using the two different conflict resolution mechanisms; naive analysis and truth discovery analysis. Two Venn diagrams in Fig. 10 show how the values from the table are related for both the naive analysis and the truth discovery analysis using the securiCAD claims as an example.

The row “Initial claims” in Table 8 shows the number of all original claims included in the analysis. “Generated negative” shows the amount of the negative claims added to the truth discovery analysis to factor in uncertainty. “Conflicting claims” shows the total number of conflicting claims that need to be analyzed. “Non-conflicting claims” shows the amount of claims that are not included in any conflict resolution, because this subset contains no conflicts. The number of conflicts solved with either of the algorithms is shown as “Conflicts solved claims”. The total number of elements that can be used to create the model, including implicit and explicit associations, is shown as “Final non-negative claims”. For the truth discovery analysis, this means that the negative claims that are found to have high credibility values are excluded together with the claims they negate.

The conflict resolution results for the three studied models are different, not only because of the two different conflict resolution methods used, but

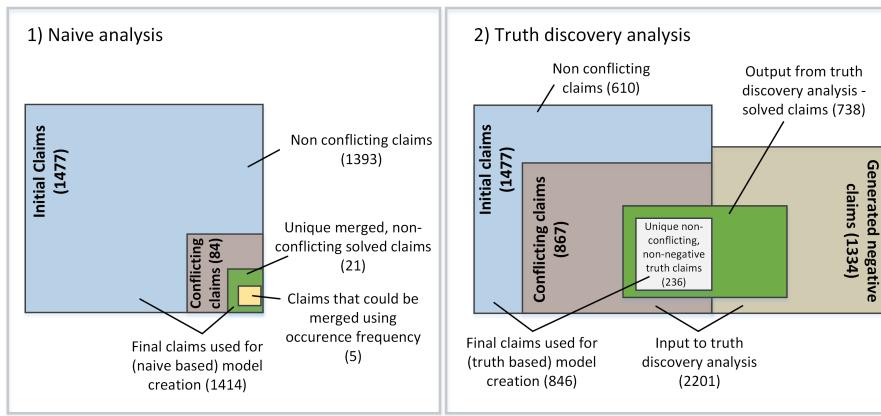


Fig. 10 Venn diagrams of securiCAD claims in naive and truth discovery analysis.

Table 9 Number of different model objects created using naive and truth analysis in the securiCAD modeling case.

| Class name | Non-merged | Naive analysis | Truth analysis |
|-----------------|------------|----------------|----------------|
| AccessControl | 6 | 4 | 4 |
| Client | 130 | 130 | 125 |
| Dataflow | 122 | 122 | 62 |
| Host | 94 | 46 | 24 |
| Network | 41 | 20 | 2 |
| Protocol | 122 | 122 | 62 |
| Router | 5 | 5 | 3 |
| Service | 681 | 661 | 252 |
| SoftwareProduct | 88 | 26 | 21 |
| UserAccount | 67 | 66 | 66 |

Table 10 Number of different model objects created using naive and truth analysis for the ArchiMate case.

| Class name | Non-merged | Naive analysis | Truth analysis |
|----------------------|------------|----------------|----------------|
| CommunicationPath | 122 | 122 | 92 |
| Device | 99 | 46 | 39 |
| Network | 41 | 20 | 20 |
| SystemSoftware | 811 | 791 | 329 |
| SystemSoftwareUnique | 88 | 26 | 18 |

Table 11 Number of different model objects created using naive and truth analysis for the EMFTA case.

| Class name | Non-merged | Naive analysis | Truth analysis |
|------------------|------------|----------------|----------------|
| interface | 99 | 46 | 15 |
| node | 94 | 46 | 14 |
| softwareInstance | 811 | 791 | 289 |
| zone | 41 | 20 | 3 |

also because of the difference in the number of included model elements (see Tables 9, 10, 11) and the nature of the associations between those elements. The truth discovery based conflict resolution could be run on all three claim sets while the naive could be run only on securiCAD and ArchiMate. The naive method could not be used for EMFTA because no conflicts in the data are identified due to the nature of associations in the model. The problem is that the metamodel specification does not restrict associations for the elements that are included in the model creation, so all the associations of the captured elements are one-to-many and thus there are no conflicts in the data. According to the specification, everything captured can exist.

Starting with the naive analysis, the conflicting data sets of securiCAD and ArchiMate include mostly operating system names. The naive analysis of the two data sets shows that very few conflicts can be solved by simply counting occurrence frequencies. For securiCAD model there are 84 conflicts for 21 claims. For ArchiMate there are 82 conflicts for 20 claims. For both models, only 5 claims could be decided using occurrence frequency. A simplistic approach to calculate trustworthiness scores for the involved data sources helped to solve the rest of the naive conflicts. The third columns in tables 9 and 10 show the count of created elements using the naive analysis. No uncertainty is factored in and thus the data sets are very similar to the original data. This creates some problems as Wireshark for example claims 20 different networks, 16 of which are not claimed by any other data source. Wireshark sees those networks because the SCADA software on the servers is trying to connect (surprisingly) to some previously configured internet addresses. The number of application software programs in the model is also inaccurate because different data sources have different naming conventions and therefore some software is represented multiple times.

Unlike the naive analysis, the truth discovery analysis uses probabilistic calculations. The results from the trustworthiness analysis of the three models

Table 12 The table shows data source trustworthiness scores (“Trust”), number of claims in the trust analysis (“Claims”) and the amount of negative values added (“Negs”).

| Data source name | securiCAD | | | ArchiMate | | | EMFTA | | |
|----------------------|-----------|--------|------|-----------|--------|------|-------|--------|------|
| | Trust | Claims | Negs | Trust | Claims | Negs | Trust | Claims | Negs |
| Nessus | 0.35 | 262 | 310 | 0.36 | 262 | 310 | 0.48 | 258 | 313 |
| Nexpose | 0.60 | 125 | 484 | 0.62 | 125 | 484 | 0.81 | 84 | 487 |
| Nmap | 0.86 | 203 | 369 | 0.82 | 203 | 369 | 0.67 | 204 | 367 |
| p0f Sec. serv. | 0.97 | 150 | 27 | 0 | 120 | 25 | 0.98 | 61 | 21 |
| Wireshark Sec. serv. | 0.02 | 30 | 144 | 0.99 | 28 | 114 | 0.06 | 27 | 61 |
| etc-passwd SCADA A | 1.0 | 1 | | 1.0 | 1 | | 1.0 | 1 | |
| etc-passwd SCADA B | 1.0 | 1 | | 1.0 | 1 | | 1.0 | 1 | |
| HP-Procurve | 1.0 | 1 | | 0.99 | 2 | | 0.99 | 2 | |
| netstat SCADA A | 0.02 | 50 | | 0.02 | 50 | | 0.02 | 50 | |
| netstat SCADA B | 0.03 | 33 | | 0.03 | 33 | | 0.03 | 33 | |
| pfSense_Config | 0.4 | 3 | | 0.71 | 5 | | 0.76 | 5 | |
| PS-Get-User AD | 1.0 | 2 | | 1.0 | 1 | | 1.0 | 1 | |
| PS-Get-User DE | 1.0 | 2 | | 1.0 | 2 | | 1.0 | 3 | |
| PS-Win32_Prod. AD. | 1.0 | 2 | | 1.0 | 2 | | 1.0 | 3 | |
| PS-Win32_Prod. DE | 1.0 | 1 | | 1.0 | 1 | | 1.0 | 1 | |

are claim credibility and data source trustworthiness scores. Trustworthiness scores for the models are shown in Table 12, and are explained in the following paragraphs.

Table 12 shows the trustworthiness scores for three types of data sources; (i) for network scanners, (ii) for network traffic capturing and filtering tools and (iii) for data sources for individual systems.

Out of the three types, the network scanners provide the easiest way to gather wide scale system information. The scanners have the largest number of claims about the IT architecture as compared to other data sources and also the largest number of negative claims generated for truth discovery. The number of negative claims is large because the framework is configured to treat the three scanners as equally capable if network zone and class restrictions are met. All three scanners pick up network zone, host and software information. However, out of the three Nessus has the lowest scores for all three models (0.35/0.36/0.48). Looking at the data sets, it is apparent that the data from Nmap and Nexpose are more similar to each other in terms of conventions for naming services and operating systems than to Nessus. Nexpose’s trustworthiness score is highest for the EMFTA analysis (0.81) very likely because the EMFTA model mostly contains network information (zones, nodes, network

interfaces) and Nmap detected slightly more network zones and hosts than did Nessus and Nmap.

Wireshark and p0f can capture network traffic between hosts and p0f in some cases is also able to detect the version of an operating system running on a detected host. p0f has higher trust scores (0.97 and 0.98) for securiCAD and EMFTA model than Wireshark (0.02 and 0.06), but for the ArchiMate model the scores are the other way around. Wireshark has high trust score (0.99), while p0f is assigned a very low trust score (near 0). The difference can be explained by the differences between models and particularly by how the different models use different subsets of the data sets generated by p0f and Wireshark, leading to different conflict patterns in each models. Thus, in addition to the network and host information the securiCAD model also uses information from p0f and Wireshark about data flows and protocols, whereas Archimate only uses information about data flows. In the EMFTA model, however, data flows, protocols and related associations are not used at all. This difference in information usage is also evident from the amount of conflicting claims for the two data sources.

Configuration data for individual systems are useful for the creation of the models because of their high accuracy. This is also evident in the truth discovery analysis. Windows and Linux configuration sources provide information such as the names of the installed software and open network ports and information about user accounts. The analysis of the Windows and Linux configuration data shows that few other sources in the lab are able to provide such detailed information about specific functionalities in a single system. This detailed information, which also has limited scope, thus causes little conflict in our setting. PowerShell cmdlets and etc-passwd as data sources get high trust scores for all the models in the analysis (1.0). Netstat is the only exception because it is the source of information about services running on hosts. Some of the service information it provides conflicts with the service information from Nessus, Nmap and Nmap. The conflicts are in many cases due to differences in naming conventions. This difference causes the lower scores for all three models.

There are two sources of network configuration data; a pfSense firewall and an HP Procurve Switch. Both provide information about the network configuration of the lab, including VLANs, but also device settings like name and network address. The HP switch provides largely unique information and therefore almost none of its data ends up as conflicting, entailing high scores

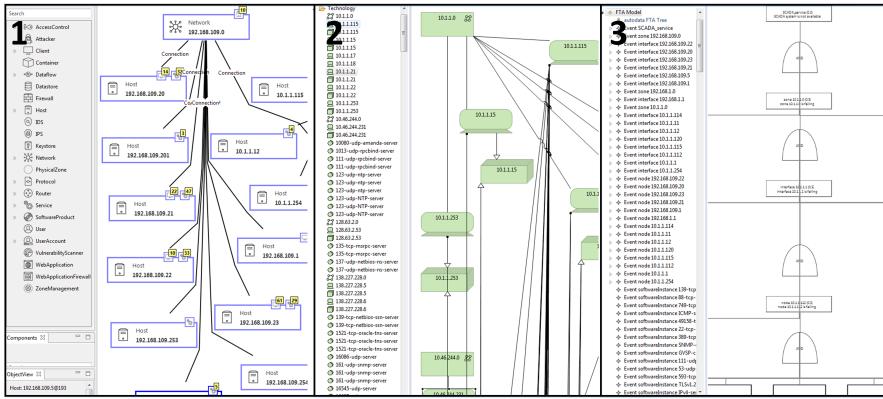


Fig. 11 The securiCAD (1), Archi (2) and EMFTA (3) models as created using the framework.

(0.99). As pfSense provides similar information to other sources like Wireshark and p0f and the network scanners, there are more conflicts and these result in a lower score.

A concluding remark on the truth analysis is that the number of elements created for each of the three models differ because, due to the model specification differences, the conflicts from the data sources vary.

Most of the configuration files are taken from the hosts in two network zones. Network scanners Nmap and Nessus are also only able to scan hosts on those two network zones (due to a firewall misconfiguration). The network aspect of securiCAD and EMFTA results are similar to the actual lab setup, where the most common hosts on the best known network zones are detected. It is more difficult to judge the open ports on the hosts as the tools used compete with each other and no ground truth is known. Once the analysis process has been successfully finished, it is time for the model import part. The purpose of this process is to import the merged data into the modeling tools. In all cases the import file generator needs to write an XML file that can then be opened as a model. Fig. 11 shows selected parts of the three imported models, since (for obvious reasons) the entire models are too big to visualize in one or even three figures.

5 Discussion and Future Work

The following subsections discuss the work in relation to the research questions, effort required for using the framework and the metamodels' coverage, choice of the truth algorithm and future work.

5.1 Extensible automated modeling using heterogeneous data sources

We raised two research questions: The first goal was to investigate how automated modelling can use multiple heterogeneous data sources in an extensible manner. The second one was how to keep track of the trustworthiness of the data sources when merging data. As a result of a literature review and extensive testing, a framework is proposed and a study with three metamodels is conducted to demonstrate the viability of the approach. This subsection discusses the two research questions raised in the introduction.

How can automated system modelling use multiple and heterogeneous sources in an extensible way?

A previous study showed that it is possible and preferable to collect data for automatic modeling [21]. The authors use data from one single network scanner to generate an ArchiMate model. However, there are other useful data sources that can be employed for model generation. The possibility to merge data from multiple sources is not explored in the previous work. This is instead the focus of our current study. It leads us to the question of the extensibility of the approach, which we define as the support for multiple changing heterogeneous data sources and the support of multiple metamodels.

In our study there are a variety of data sources used that have different degrees of granularity and credibility. There are firewall and switch configuration files that describe particular devices in an idealized way (assuming that a device works correctly) and then there is information from network scanners. The methods that network scanners use to obtain data can in some cases be considered less trustworthy (like fingerprinting) and also disruptive. Obtaining individual configuration files takes more time and effort than setting up network scanners, but configuration files can contain more precise information.

The metamodels considered in this paper are very different in terms of complexity. For example, securiCAD is a graph based approach that has a very elaborate structure with many classes and rules for how these relate. EMFTA, on the other hand, has only a handful of classes used to build a

tree structure that can be applied in different contexts. To use EMFTA for system modeling one also needs to define the dependencies between all the systems. Given the varying complexity, all three metamodels have different data needs, but it is difficult to see how any one single data source would suffice for instantiating any of them. Moreover, different models might be required for enterprise decision support. Therefore, it can be concluded that by supporting model generation of multiple models based on multiple heterogeneous data sources, the authors have improved the previous approach.

The quality of the output models depends on the quality of the input data. Heterogeneous data comes in different representations and coverage. Data from different sources can be conflicting and these conflicts need to be resolved during a merging process. For conflict resolution it is important that data from different sources is comparable. Therefore, steps need to be taken to make sure that the same properties are used for describing the same elements and that data are on the same formats.

In order to make the collected data comparable, it is transformed into sets of dependent claims. These claims are eventually transformed into long textual strings (see 3.6.3) so that they can be compared with each other and so that their occurrences can be counted. This method is successfully applied to all three metamodels in the study. However, only fairly simple data warehousing techniques are used to clean and standardize the data and this affects the results. There are likely multiple representations in the data for the same object. Another issue is that it is difficult to classify some model elements, such as software. Both of these problems need external help to be solved, e.g. using an ontology. Some missing values in the data are approximated. One of the reasons for this is that the data sets are stored as a graph in the database. Other reasons are the dependencies of model classes and the need to construct long strings for comparison. For example, if a scanner like Wireshark says that there is a data flow then it is also assumed that there is a network(s) where this data flow travels and hosts with software allowing this flow.

The database facilitates consolidating similar data and tracing back original claims. Tracing back original claims becomes possible by storing data sets after each transformation and using the same database-assigned identification numbers for claims at each step of the process. These identification numbers are needed to support data provenance. The identification numbers are kept even after the conflict resolution and data merging. The identical claims from

different sources are grouped and each group is assigned a new identification number and the source info is discarded.

In conclusion, we can say that the tests with the framework show that heterogeneous data can be merged in an extensible way to create multiple models. The data sets first need to be transformed to a common representation and granularity level and stored in a database for further analysis. Then resolution of conflicts in the data becomes possible.

How can trustworthiness of multiple and heterogeneous sources be included in automated system modelling?

Data warehouses are often used for integrating distributed information sources. A data warehouse is very labor intensive to set up and to our knowledge it is difficult to track (and express) the trustworthiness of different data sources. Truth discovery is a mature field and its probabilistic methods, which mostly have been applied in a World Wide Web context, can also be used in an enterprise to automate modeling. For the sake of comparison, the authors developed an additional simple non-probabilistic method, which we call the naive approach, to track the trustworthiness of data sources. The naive approach counts occurrences of identical textual values (strings) and uses those counts to calculate scores for data sources. The idea for developing the naive approach is to evaluate the usefulness of truth discovery as compared to more conventional frequency analysis.

The results show that a naive approach using frequency based analysis for conflict resolution has limited use. The naive approach can only handle mutually exclusive claims and cannot express the reliability problems a data source could have. In the securiCAD and ArchiMate cases only approximately 20% of the conflicts can be resolved with simple frequency analysis, where the most frequent value is identified. In the remaining cases the naive algorithm has to resort to alternative methods. A probabilistic approach is preferable if there is doubt about the trustworthiness of data sources. A truth algorithm, such as LCA, is able to calculate the trustworthiness of data sources and the credibility of claims even if there are only a few frequent values. The truth discovery algorithm that was chosen calculates the probability that a data source asserts the correct claim, while prioritizing claims from the data sources that have the most matching claims.

Another important benefit of using the truth discovery algorithm over the naive one is that it allows to use prior values and factor in uncertainty with negative claims. As a result of using negative claims, some data get excluded

from the final results. These are the claims that originate from the data sources that have been attributed the lowest trustworthiness values and that have the highest number of conflicts that are removed. The extent to which data gets excluded can be regulated with negative claim generation and prior values that describe the user's trust in the data sources.

In summary, we can say that a probabilistic approach is preferable if the goal of the model generation is to have an abstract model with less disputed claims. On the other hand, if the goal is to keep all the information and not to take the probabilistic characteristics such as trustworthiness of the data sources and the credibility of the claims into account in the model creation, then a simpler approach would work as well.

5.2 Effort and coverage

The results show that when setting up the framework for usage with the three metamodels time and effort is needed to configure it. However, given that about a thousand model elements are created, manual modeling would have taken much longer. Since frequent updating is beneficial, most likely required, manual modeling also becomes unfeasible from a life-cycle perspective. A comparison of how much time manual and automated modeling can take has been done by Holm et al.[21]. Regarding the time spent setting up the framework, most of it was spent creating mappings between different model languages, including the mapping between data source elements and the common language. Correct mappings between the different data and metamodel elements are vital to obtain an accurate holistic systems and software (enterprise) model. Other issues that took time include finding the right abstraction level for the input data and common representation of data values. For example, only the network traffic that captures values that described application protocols from Wireshark are imported and the protocol port numbers are replaced with known application protocol names.

Although our study tries to improve the coverage of the data for the three chosen metamodels, this is only partly achieved due to multiple factors: 1) Limited to the devices and services running in the studied SCADA environment, some types of devices and software like an IDS or IPS are simply not present. 2) No business process data are included in the analysis, which means that the ArchiMate model is only partially covered. If there would have been a data source for business processes then the framework would have allowed

it to be added. Johnson et al. [23] discuss possible data sources for the other layers. And 3) partly, the incompleteness of the models is caused by the inability to classify software and hardware products and services. For example we cannot identify what software products are firewalls, anti-viruses, data stores, or web applications. This type of classification could be achieved with external help (e.g. ontologies). For the securiCAD model, there are enough data available to create the structure of the systems and software (IT) architecture, but there is information lacking about many of the properties of the model objects. Some of this information can be collected from the systems using other tools, however some would require external help or manual work. An example where additional tools would be useful are the properties of a *Host*, namely “ASLR”, “DEP”, “StaticARPTables”, “AntiMalware”, and “HostFirewall”. These could theoretically be collected automatically with the appropriate tools. E.g. Windows 7 “ASLR” extension can be identified with Process Explorer³³. One example of where external help is needed are some of the properties of the class *SoftwareProduct*, e.g. “StaticCodeAnalysis”, “SecretBinary”, “SecretSource”, and “HasVendorSupport”. None of these can be obtained from the data collected from the systems. This information could be collected manually from the product vendors or development organization instead. External help is also needed when it comes to classifying the instances of *SoftwareProducts*. For example, one needs to identify which software is client software, server software, web servers, data stores, vulnerability scanners, and anti-malware.

For the EMFTA model it becomes obvious that the probability values for the faults in the fault tree elements cannot be collected automatically from our lab. Thus, the generated EMFTA model might seem incomplete in that it lacks quantitative data and therefore cannot be used for reliability calculations. However, even though precise figures are lacking, approximate calculations can nevertheless be run. For the particular kind of system used in our case study – a SCADA system for electric power system control and operation – availability figures for all the ICT components can be found in literature e.g. [22]. Even though the granularities of that model and the one presented in this paper do not match perfectly, it would nevertheless be possible to translate the figures so that they fit the model presented as a first approximation. Another route to approximate calculations would be to use distributions from empirical studies

³³ <https://blog.didierstevens.com/2011/01/18/quickpost-checking-aslr>

of ICT components, such as Weibull of time-to-failure [44] [19] and log-normal of time-to-recovery [44] [14]. If generic estimates are not deemed good enough, it is also possible to elicit the relevant numbers, i.e. time-to-failure and time-to-recovery, for the system at hand using stakeholder interviews. This has been shown to be feasible in practice without too much effort [32].

An attempt was made to complement the data with outside source information to improve the quality when importing data with adapters. An example of such a complement is the IANA’s Service Name and Transport Protocol Port Number Registry³⁴, which add service names to p0f traffic captures. Another possibility is to organize Linux packages into application type based groups. Some Linux vendors provide lists of packages arranged according to their purpose. Such lists are available for example for some versions of Ubuntu³⁵ and can be used to classify the detected software according to model needs (core packages, servers, clients etc).

5.3 Truth algorithm and negative values

There are many different truth discovery algorithms available based on various assumptions and characteristics. The main purpose of many of them is to reason over the truthfulness of claims made by different web sources. Some algorithms assume that there are sources that lie on purpose, while others factor in accidental errors and include a degree of randomness. The authors of this paper chose one algorithm, LCA, to reason over system data based on tests of multiple algorithms. The algorithm chosen is easy to understand and tractable. It has fewer parameters than other similar algorithms thus the possibility of overfitting is smaller. While some effort was spent on different assumptions of the truth discovery models and picking an effective algorithm, the authors do not claim that the algorithm chosen is the most suitable one for reasoning over IT architecture data. The purpose of this work is to show that using a truth discovery algorithm for solving heterogeneous problems in IT architecture data makes sense. Picking the best algorithm needs a separate study, which focuses mainly on algorithm comparison.

Input errors might occur especially when it comes to using manually created data. Simpler and common errors are fixed with cleaning and filtering

³⁴ <https://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml>

³⁵ <https://packages.ubuntu.com>

techniques implemented in the data source specific adapters, while more complex ones require external help, e.g. by using ontologies. Data errors that are not resolved will end up as conflicts in the truth discovery analysis. Given that there are enough data sources and one data source keeps making errors, these errors will reduce the trustworthiness of the data source that thus will be assigned a low credibility score. It will eventually be excluded from the final data set by the truth discovery analysis.

As explained, the trustworthiness scores are influenced by negative claims. For that reason the generation of negative claims should be considered carefully for each data source. It can be seen from the results that some data sources such as network scanners have extensive coverage while others like PowerShell cmdlets only made claims about single systems. The potential capability to have information about particular services and network zones is also a reason for the high number of negative values that are generated for some sources. E.g. for Nexpose the number of added negative claims is higher than the number of original claims. The negative values influence the trustworthiness score of that particular data source, but also the trustworthiness scores of other data sources whose claims are being disputed.

5.4 Summary and future work

The authors believe that the presented study shows that the proposed framework has practical value for IT architecture modeling using heterogeneous data sources to automate the process. One probabilistic truth discovery algorithm is tested and it shows promising results by filtering out highly disputed data values. The strength of the presented framework lies within its modularity. Different algorithms can be plugged in for conflict resolution as demonstrated. The study also shows that external help such as an ontology is needed in order to fully make use of the data. If data values can be classified and grouped in a meaningful way, the final models will be even more informative.

Future work needs to focus on improving the quality of the input data. That would help increasing the quality of the output models and reducing the time needed for data processing. Other potential topics include a comparative study of truth discovery models and assessing the framework by applying it in other environments with more diverse data sources.

6 Conclusion

This paper proposes a framework for automatically creating holistic IT architecture models from machine-readable data. The framework supports heterogeneous input data to improve the coverage of different metamodels. Data are merged in a way where a truth algorithm reasons over the trustworthiness of different sources. The calculations take into account the trust in the sources and the claims each source is making. As a result, highly disputed data are excluded from the final models.

The framework is tested on a Supervisory Control And Data Acquisition (SCADA) environment. Three models are automatically created in order to show the extensibility of the framework. The calculation results of the credibility values are compared to a simpler alternative – naive frequency analysis. Our tests suggest that the proposed framework can be used to speed up IT architecture modeling with automation when heterogeneous data from multiple sources is used for the generation of more than one type of IT architecture model.

Acknowledgements This work has received funding from the EU FP7 under grant agreement no. 607109 (SEGRID) and the Swedish National Grid.

References

1. Aier, S., Buckl, S., Franke, U., Gleichauf, B., Johnson, P., Naerman, P., Schweda, C.M., Ullberg, J.: A survival analysis of application life spans based on enterprise architecture models. In: J. Mendling, S. Rinderle-Ma, W. Esswein (eds.) Enterprise modelling and information systems architectures : proceedings of the 3rd International Workshop on Enterprise Modelling and Information Systems Architectures, vol. LNI P-152, pp. 141–154. Ges. für Informatik, Bonn (2009). URL <https://www.alexandria.unisg.ch/213537/>
2. Alegria, A., Vasconcelos, A.: IT architecture automatic verification: A network evidence-based approach. In: Research Challenges in Information Science (RCIS), 2010 Fourth International Conference on, pp. 1–12. IEEE (2010)
3. Berti-Equille, L., Borge-Holthoefer, J.: Veracity of data: From truth discovery computation algorithms to models of misinformation dynamics. *Synthesis Lectures on Data Management* **7**(3), 1–155 (2015)
4. Borlase, S.: Smart Grids: Infrastructure, Technology, and Solutions. CRC press (2016)
5. Buschle, M., Holm, H., Sommestad, T., Ekstedt, M., Shahzad, K.: A tool for automatic enterprise architecture modeling. In: Forum at the Conference on Advanced Information Systems Engineering (CAiSE), pp. 1–15. Springer (2011)
6. Cheney, J., Chiticariu, L., Tan, W.C.: Provenance in databases: Why, how, and where. *Foundations and Trends in Databases* **1**(4), 379–474 (2009). DOI 10.1561/1900000006. URL <http://dx.doi.org/10.1561/1900000006>

7. Dayal, U., Castellanos, M., Simitsis, A., Wilkinson, K.: Data integration flows for business intelligence. In: Proceedings of the 12th International Conference on Extending Database Technology: Advances in Database Technology, EDBT '09, pp. 1–11. ACM, New York, NY, USA (2009). DOI 10.1145/1516360.1516362. URL <http://doi.acm.org/10.1145/1516360.1516362>
8. El-Sappagh, S.H.A., Hendawi, A.M.A., Bastawissy, A.H.E.: A proposed model for data warehouse etl processes. Journal of King Saud University - Computer and Information Sciences **23**(2), 91 – 104 (2011). DOI <https://doi.org/10.1016/j.jksuci.2011.05.005>. URL <http://www.sciencedirect.com/science/article/pii/S131915781100019X>
9. Farwick, M., Greiter, B., Breu, R., Ryll, S., Voges, K., Hanschke, I.: Requirements for automated enterprise architecture model maintenance. In: 13th International Conference on Enterprise Information Systems (ICEIS), Beijing (2011)
10. Farwick, M., Breu, R., Hauder, M., Roth, S., Matthes, F.: Enterprise architecture documentation: Empirical analysis of information sources for automation. In: System Sciences (HICSS), 2013 46th Hawaii International Conference on, pp. 3868–3877. IEEE (2013)
11. Farwick, M., Schweda, C.M., Breu, R., Hanschke, I.: A situational method for semi-automated enterprise architecture documentation. Software & Systems Modeling **15**(2), 397–426 (2016)
12. Frank, U.: Enterprise modelling: The next steps. Enterprise Modelling and Information Systems Architectures—International Journal of Conceptual Modeling **9**(1), 22–37 (2014)
13. Franke, U., Ekstedt, M., Lagerström, R., Saat, J., Winter, R.: Trends in enterprise architecture practice—a survey. Trends in Enterprise Architecture Research pp. 16–29 (2010)
14. Franke, U., Holm, H., König, J.: The distribution of time to recovery of enterprise IT services. IEEE Transactions on Reliability **63**(4), 858–867 (2014)
15. Glavic, B., Dittrich, K.R.: Data provenance: A categorization of existing approaches. In: A. Kemper, H. Schning, T. Rose, M. Jarke, T. Seidl, C. Quix, C. Brochhaus (eds.) BTW, LNI, vol. 103, pp. 227–241. GI (2007). URL <http://dblp.uni-trier.de/db/conf/btw/btw2007.html#GlavicD07>
16. Haas, L., Miller, R., Niswonger, B., Roth, M.T., Schwarz, P., Wimmers, E.: Transforming heterogeneous data with database middleware: Beyond integration. Data Engineering **22**(1), 31–36 (1999)
17. Hacks, S., Licher, H.: Distributed enterprise architecture evolution—a roundtrip approach. In: Wirtschaftsinformatik 2017 Proceedings. AIS (2017)
18. Halevy, A., Rajaraman, A., Ordille, J.: Data integration: The teenage years. In: Proceedings of the 32Nd International Conference on Very Large Data Bases, VLDB '06, pp. 9–16. VLDB Endowment (2006). URL <http://dl.acm.org/citation.cfm?id=1182635.1164130>
19. Heath, T., Martin, R.P., Nguyen, T.D.: Improving cluster availability using workstation validation. In: ACM SIGMETRICS Performance Evaluation Review, vol. 30, pp. 217–227. ACM (2002)
20. Hinkelmann, K., Gerber, A., Karagiannis, D., Thoenissen, B., Van der Merwe, A., Woitsch, R.: A new paradigm for the continuous alignment of business and IT: Combining enterprise architecture modelling and enterprise ontology. Computers in Industry **79**, 77–86 (2016)

21. Holm, H., Buschle, M., Lagerström, R., Ekstedt, M.: Automatic data collection for enterprise architecture models. *Software & Systems Modeling* **13**(2), 825–841 (2014)
22. Jensen, M., Sel, C., Franke, U., Holm, H., Nordström, L.: Availability of a SCADA/OMS/DMS system – a case study. In: Innovative Smart Grid Technologies Conference Europe (ISGT Europe), 2010 IEEE PES, pp. 1–8. IEEE (2010)
23. Johnson, P., Ekstedt, M., Lagerstrom, R.: Automatic probabilistic enterprise it architecture modeling: A dynamic bayesian networks approach. In: Trends in Enterprise Architecture Research (TEAR), Enterprise Distributed Object Computing Workshop (EDOCW), 2016 IEEE 20th International. IEEE (2016)
24. Johnson, P., Lagerström, R., Närman, P., Simonsson, M.: Extended influence diagrams for system quality analysis. *Journal of Software* **2**(3), 30–42 (2007)
25. Jugel, D., Schweda, C.M.: Interactive functions of a cockpit for enterprise architecture planning. In: Enterprise Distributed Object Computing Conference Workshops and Demonstrations (EDOCW), 2014 IEEE 18th International, pp. 33–40. IEEE (2014)
26. Kühn, H., Bayer, F., Junginger, S., Karagiannis, D.: Enterprise model integration. *E-Commerce and Web Technologies* pp. 379–392 (2003)
27. Lagerström, R., Baldwin, C., MacCormack, A., Dreyfus, D.: Visualizing and measuring enterprise architecture: an exploratory biopharma case. In: IFIP Working Conference on The Practice of Enterprise Modeling, pp. 9–23. Springer, Berlin, Heidelberg (2013)
28. Li, Y., Gao, J., Meng, C., Li, Q., Su, L., Zhao, B., Fan, W., Han, J.: A survey on truth discovery. *Acm Sigkdd Explorations Newsletter* **17**(2), 1–16 (2016)
29. Li, Y., Li, Q., Gao, J., Su, L., Zhao, B., Fan, W., Han, J.: On the discovery of evolving truth. In: Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 675–684. ACM (2015)
30. Lozano, M.G., Franke, U., Rosell, M., Vlassov, V.: Towards automatic veracity assessment of open source information. In: 2015 IEEE International Congress on Big Data, pp. 199–206 (2015). DOI 10.1109/BigDataCongress.2015.36
31. Moser, C., Junginger, S., Brückmann, M., Schöne, K.M.: Some process patterns for enterprise architecture management. In: Software Engineering (Workshops), pp. 19–30 (2009)
32. Närman, P., Franke, U., König, J., Buschle, M., Ekstedt, M.: Enterprise architecture availability analysis using fault trees and stakeholder interviews. *Enterprise Information Systems* **8**(1), 1–25 (2014)
33. Närman, P., Johnson, P., Lagerström, R., Franke, U., Ekstedt, M.: Data collection prioritization for system quality analysis. *Electronic Notes in Theoretical Computer Science* **233**, 29–42 (2009)
34. Neaga, E.I., *, J.A.H.: An enterprise modeling and integration framework based on knowledge discovery and data mining. *International Journal of Production Research* **43**(6), 1089–1108 (2005). DOI 10.1080/00207540412331322939. URL <https://doi.org/10.1080/00207540412331322939>
35. Niemi, E., Pekkola, S.: Using enterprise architecture artefacts in an organisation. *Enterprise Information Systems* **11**(3), 313–338 (2017)
36. Nowakowski, E., Farwick, M., Trojer, T., Häusler, M., Kessler, J., Breu, R.: Enterprise architecture planning: Analyses of requirements from practice and research. In: Proceedings of the 50th Hawaii International Conference on System Sciences (HICSS) (2017)

37. Pasternack, J., Roth, D.: Latent credibility analysis. In: Proceedings of the 22nd international conference on World Wide Web, pp. 1009–1020. ACM (2013)
38. Rahm, E., Bernstein, P.A.: A survey of approaches to automatic schema matching. the VLDB Journal **10**(4), 334–350 (2001)
39. Rahm, E., Do, H.H.: Data cleaning: Problems and current approaches. IEEE Data Eng. Bull. **23**(4), 3–13 (2000)
40. Rekatsinas, T., Dong, X.L., Srivastava, D.: Characterizing and selecting fresh data sources. In: Proceedings of the 2014 ACM SIGMOD international conference on Management of data, pp. 919–930. ACM (2014)
41. Rundensteiner, E.A., Koeller, A., Zhang, X.: Maintaining data warehouses over changing information sources. Commun. ACM **43**(6), 57–62 (2000). DOI 10.1145/336460.336475. URL <http://doi.acm.org/10.1145/336460.336475>
42. Saat, J., Winter, R., Franke, U., Lagerstrom, R., Ekstedt, M.: Analysis of it/business alignment situations as a precondition for the design and engineering of situated it/business alignment solutions. In: System Sciences (HICSS), 2011 44th Hawaii International Conference on, pp. 1–9. IEEE (2011)
43. Sandkuhl, K., Stirna, J., Persson, A., Wibotzki, M.: Enterprise modeling. Springer (2014)
44. Schroeder, B., Gibson, G.: A large-scale study of failures in high-performance computing systems. IEEE Transactions on Dependable and Secure Computing **7**(4), 337–350 (2010)
45. Simmhan, Y.L., Plale, B., Gannon, D.: A survey of data provenance in e-science. SIGMOD Rec. **34**(3), 31–36 (2005). DOI 10.1145/1084805.1084812. URL <http://doi.acm.org/10.1145/1084805.1084812>
46. Trojer, T., Farwick, M., Häusler, M., Breu, R.: Living modeling of it architectures: challenges and solutions. In: Software, Services, and Systems, pp. 458–474. Springer (2015)
47. Välja, M., Lagerström, R., Ekstedt, M., Korman, M.: A requirements based approach for automating enterprise it architecture modeling using multiple data sources. In: 2015 IEEE 19th International Enterprise Distributed Object Computing Workshop, pp. 79–87 (2015). DOI 10.1109/EDOCW.2015.33
48. Veneberg, R., Iacob, M.E., van Sinderen, M.J., Bodenstaff, L.: Enterprise architecture intelligence: combining enterprise architecture and operational data. In: Enterprise Distributed Object Computing Conference (EDOC), 2014 IEEE 18th International, pp. 22–31. IEEE (2014)
49. Waguih, D.A., Berti-Equille, L.: Truth discovery algorithms: An experimental evaluation. arXiv preprint arXiv:1409.6428 (2014)
50. Wang, T., Truptil, S., Benaben, F.: An automatic model-to-model mapping and transformation methodology to serve model-based systems engineering. Inf. Syst. E-bus. Manag. **15**(2), 323–376 (2017). DOI 10.1007/s10257-016-0321-z. URL <https://doi.org/10.1007/s10257-016-0321-z>